

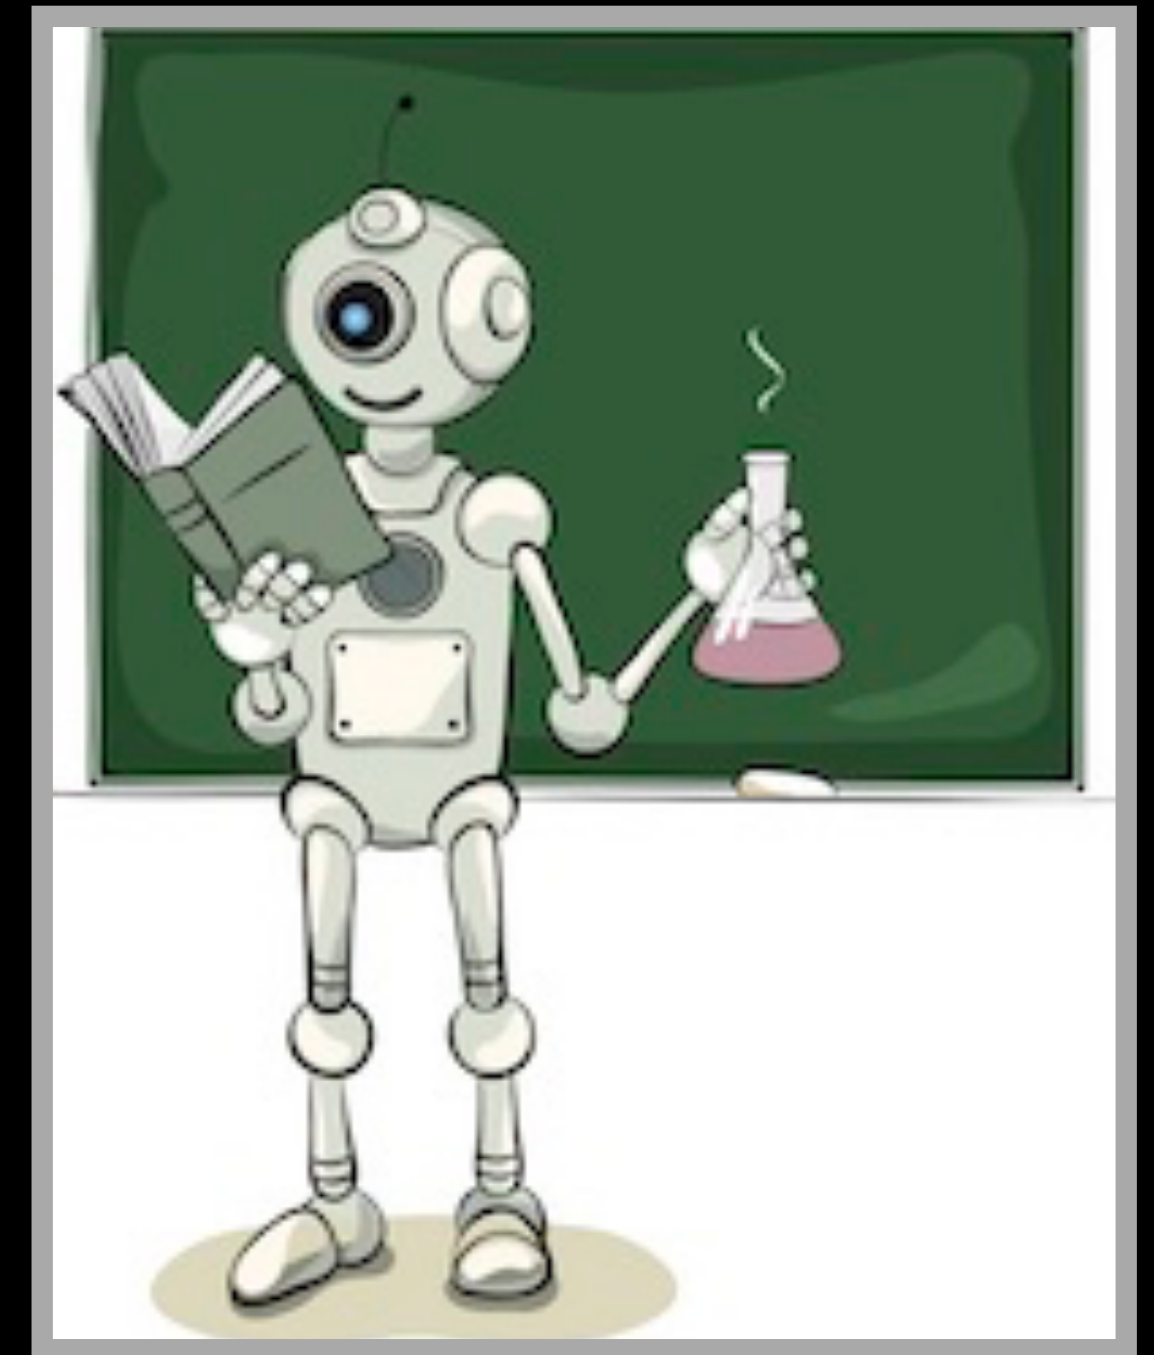
# Getting Started with Machine Learning in Python

Julian Gold, DataX Data Scientist, CSML  
Wednesday, February 28, 2024 at 4:30-6:00 PM

`https://github.com/PrincetonUniversity/python\_machine\_learning`

# What is machine learning?

A computer observes some data,  
builds a model based on the data,  
and uses the model as both a hypothesis about the world  
and a piece of software that can solve problems.



# Why use machine learning?

*Why not just program the model?*

- you may not know the model
- you want the model to be flexible and adapt as you have new data

# Why use classical machine learning?

- You want a physically interpretable model
- Your data isn't images
- You have a simple-ish problem

# Supervised Machine Learning

Computer observes input-output pairs  
and learns a function that maps input to output.

- *Labeled* data

# Supervised Machine Learning

Training set of  $N$  input-output pairs:

$$(X_1, y_1), (X_2, y_2), \dots (X_N, y_N)$$

where each pair was generated by an unknown function  $y = f(X)$ .

Goal: learn a function  $h$  that approximates the true function  $f$ .

# Unsupervised Machine Learning

Computer learns patterns/structure in input data.

- Without labels.

# Unsupervised Machine Learning

Training set of N inputs:

$$X_1, X_2, \dots, X_N$$

Goal: learn structure/patterns in the data



# Supervised Machine Learning

- You have simulated data where you know the truth
- You have data labeled by a human

# Unsupervised Machine Learning

- You have experimental data
- You want to explore / visualize the data

# Common algorithms

Supervised	Unsupervised
Regression	Clustering
Classification	Dimensionality Reduction

# Common algorithms

Supervised	Unsupervised
Regression	Clustering
Classification	Dimensionality Reduction

# Regression

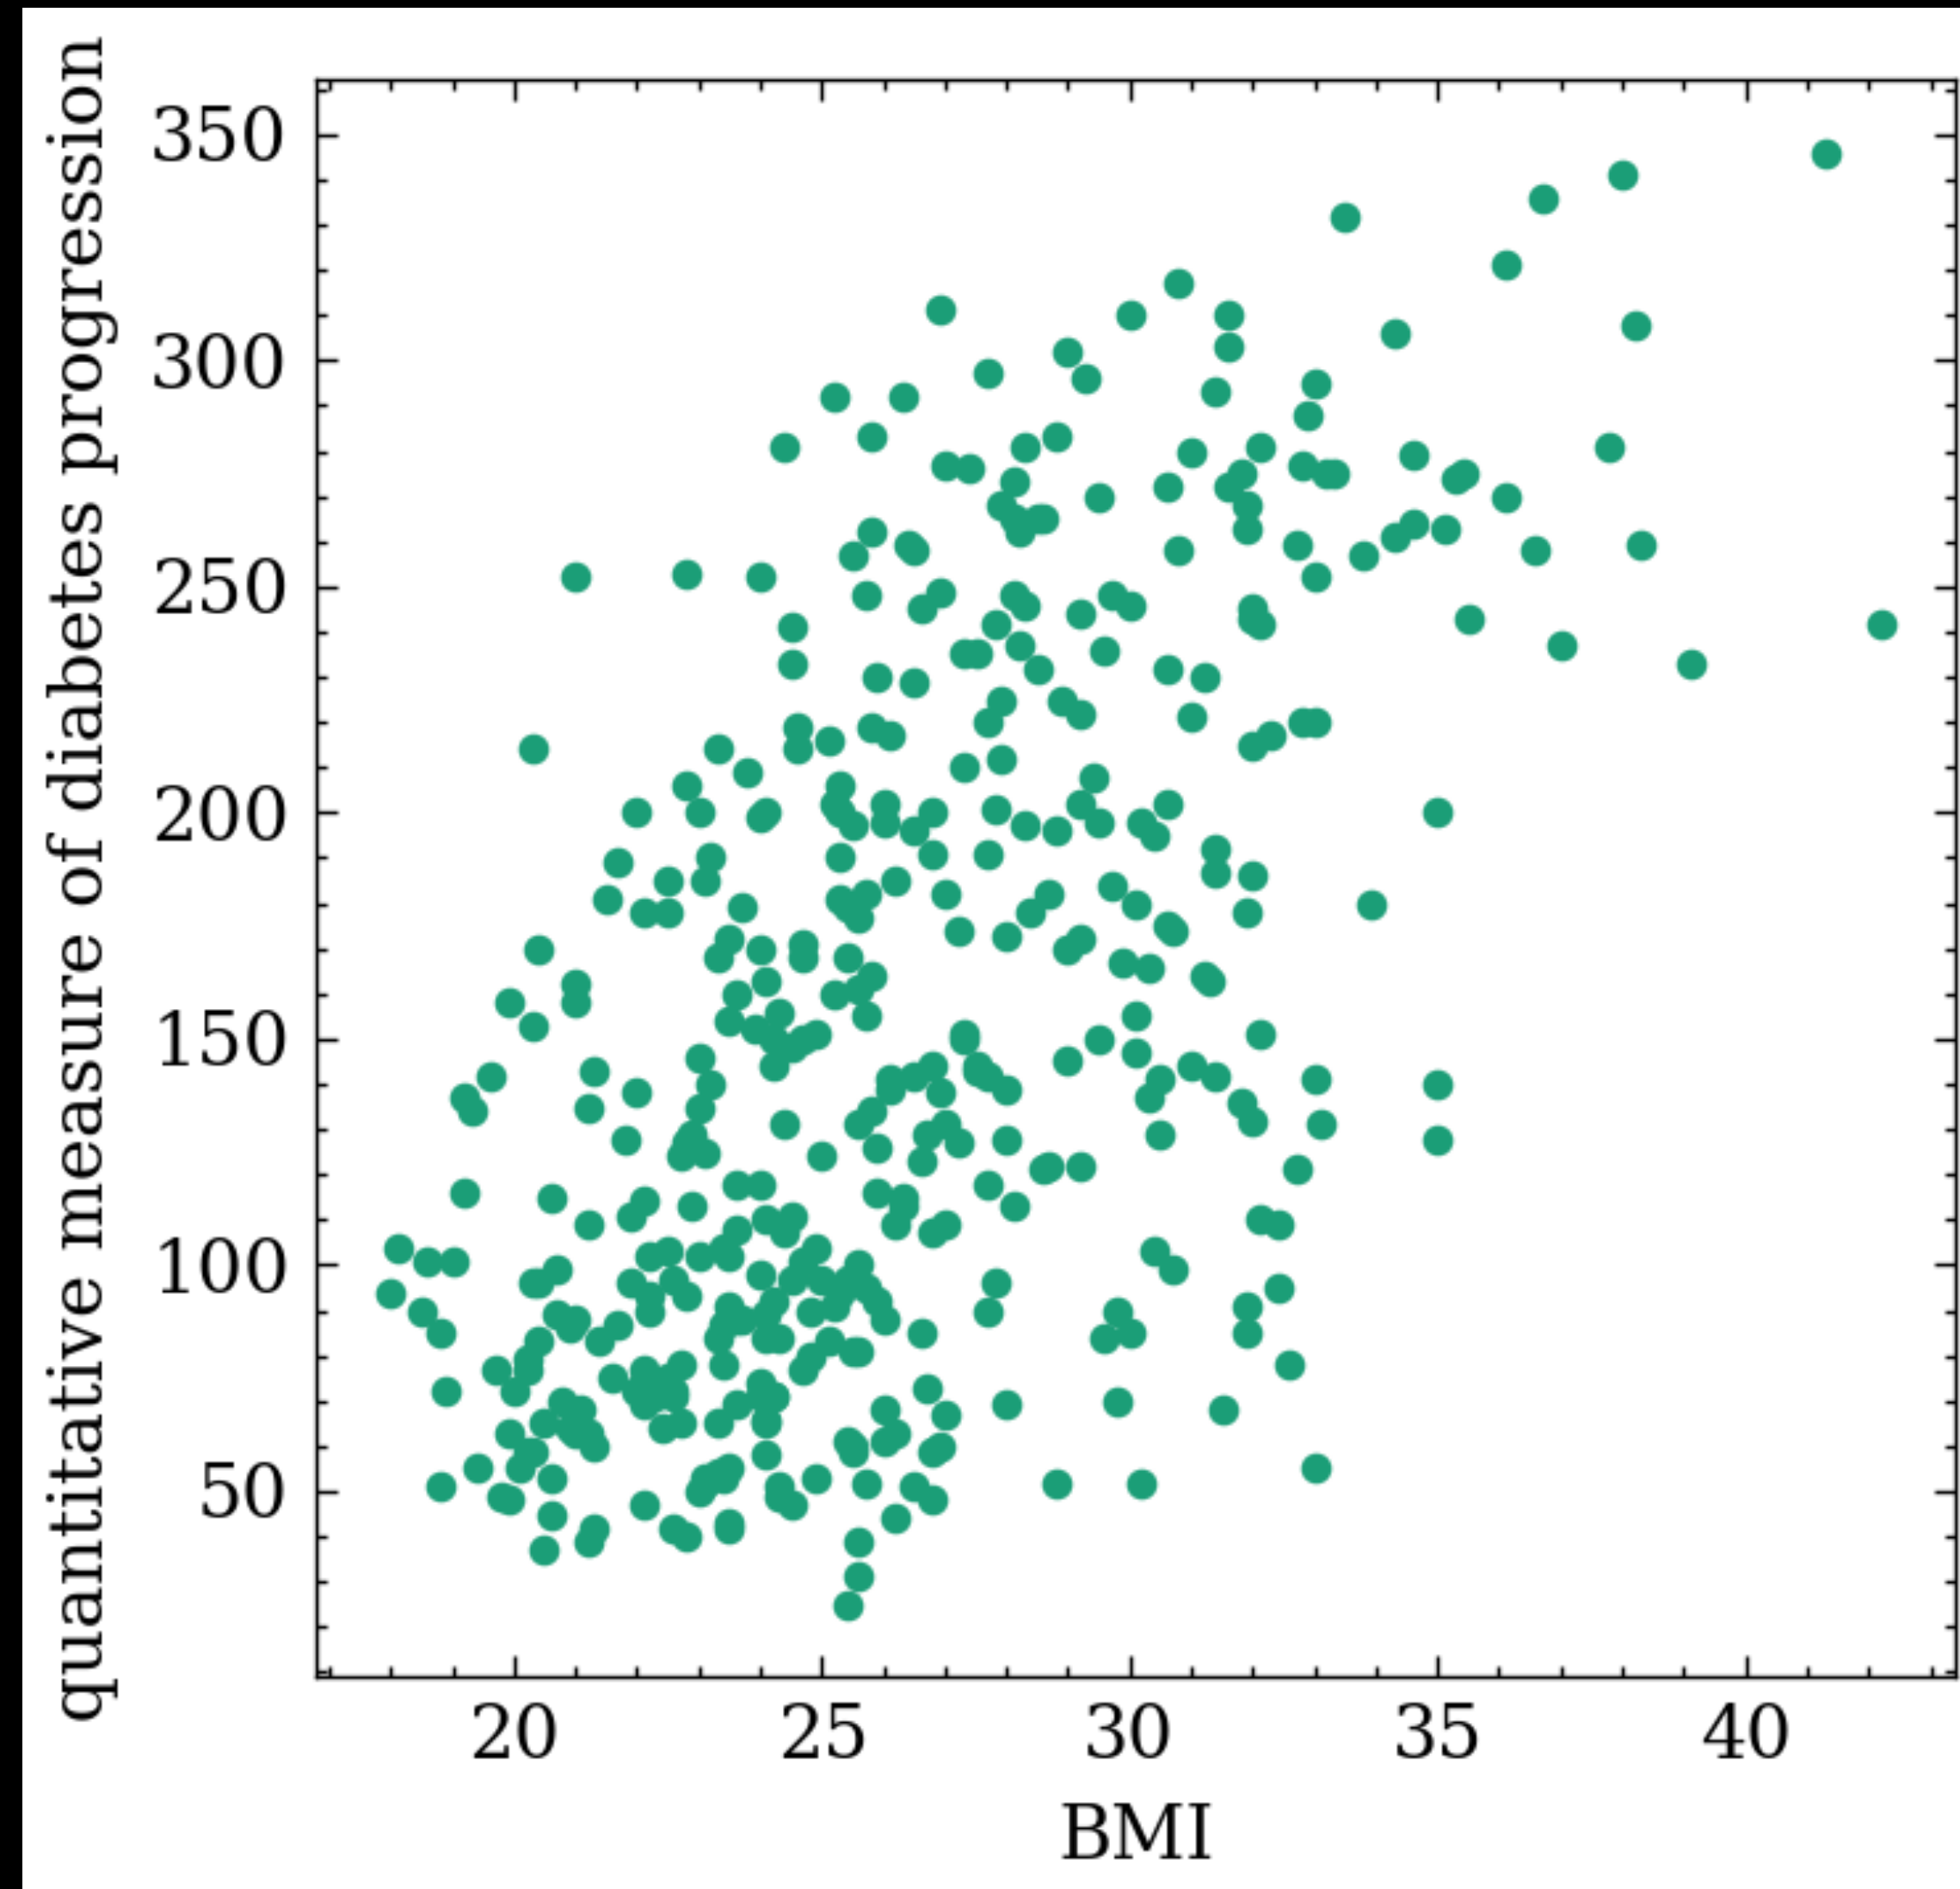
Training set of  $N$  input-output pairs:

$$(X_1, y_1), (X_2, y_2), \dots (X_N, y_N)$$

where desired output value is a continuous value.

Goal: learn a function that approximates the true function  $f$ .

# Linear Regression



input-output pairs:

- $X_i$ , BMI
- $y_i$ , diabetes progression

Diabetes dataset:

*`sklearn.datasets.load_diabetes`*

Goal: learn a function  $h$  that approximates the true function  $f$ :

*Predicted diabetes progression:  $\hat{y} = h(x)$*

# Linear Regression

$$y = b + mx$$

dependent variable

independent variables

parameter vector

$$y = X\beta$$

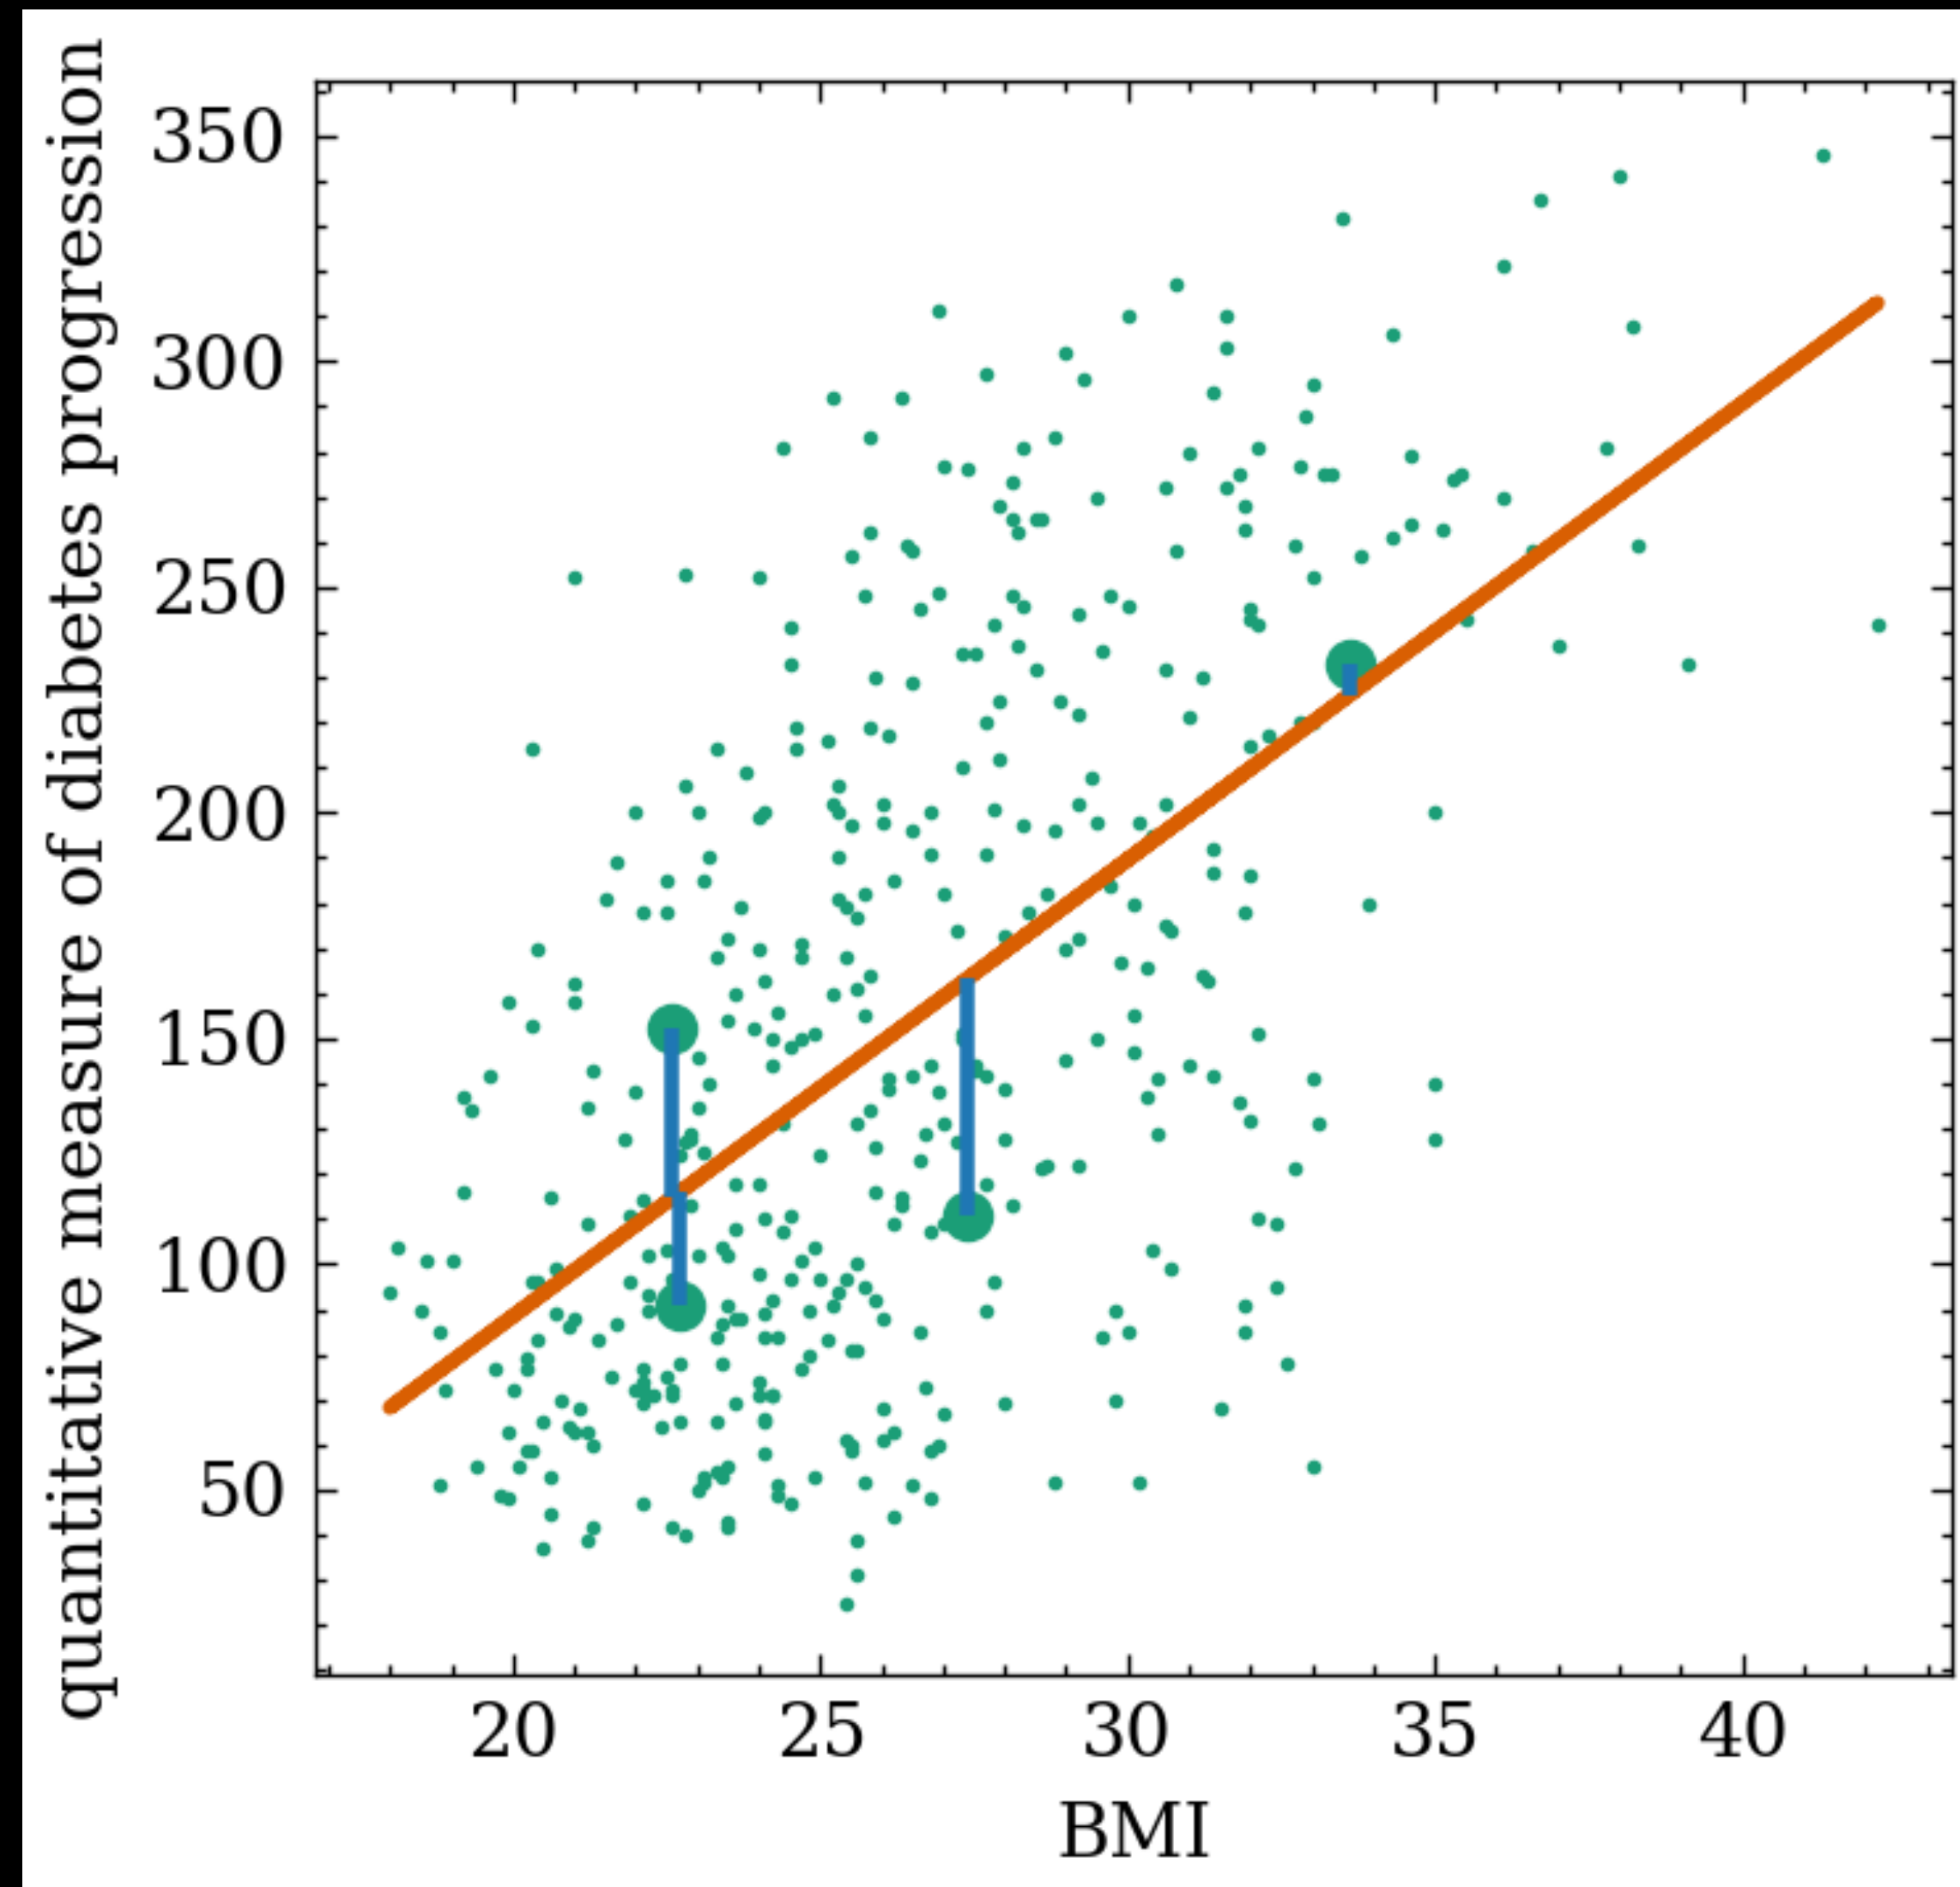
The diagram shows the matrix equation  $y = X\beta$ . An arrow points from the label 'dependent variable' to the variable  $y$ . Another arrow points from the label 'independent variables' to the matrix  $X$ . A third arrow points from the label 'parameter vector' to the vector  $\beta$ .

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$$

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix}$$

$$\beta = \begin{bmatrix} b \\ m \end{bmatrix}$$

# Linear Regression



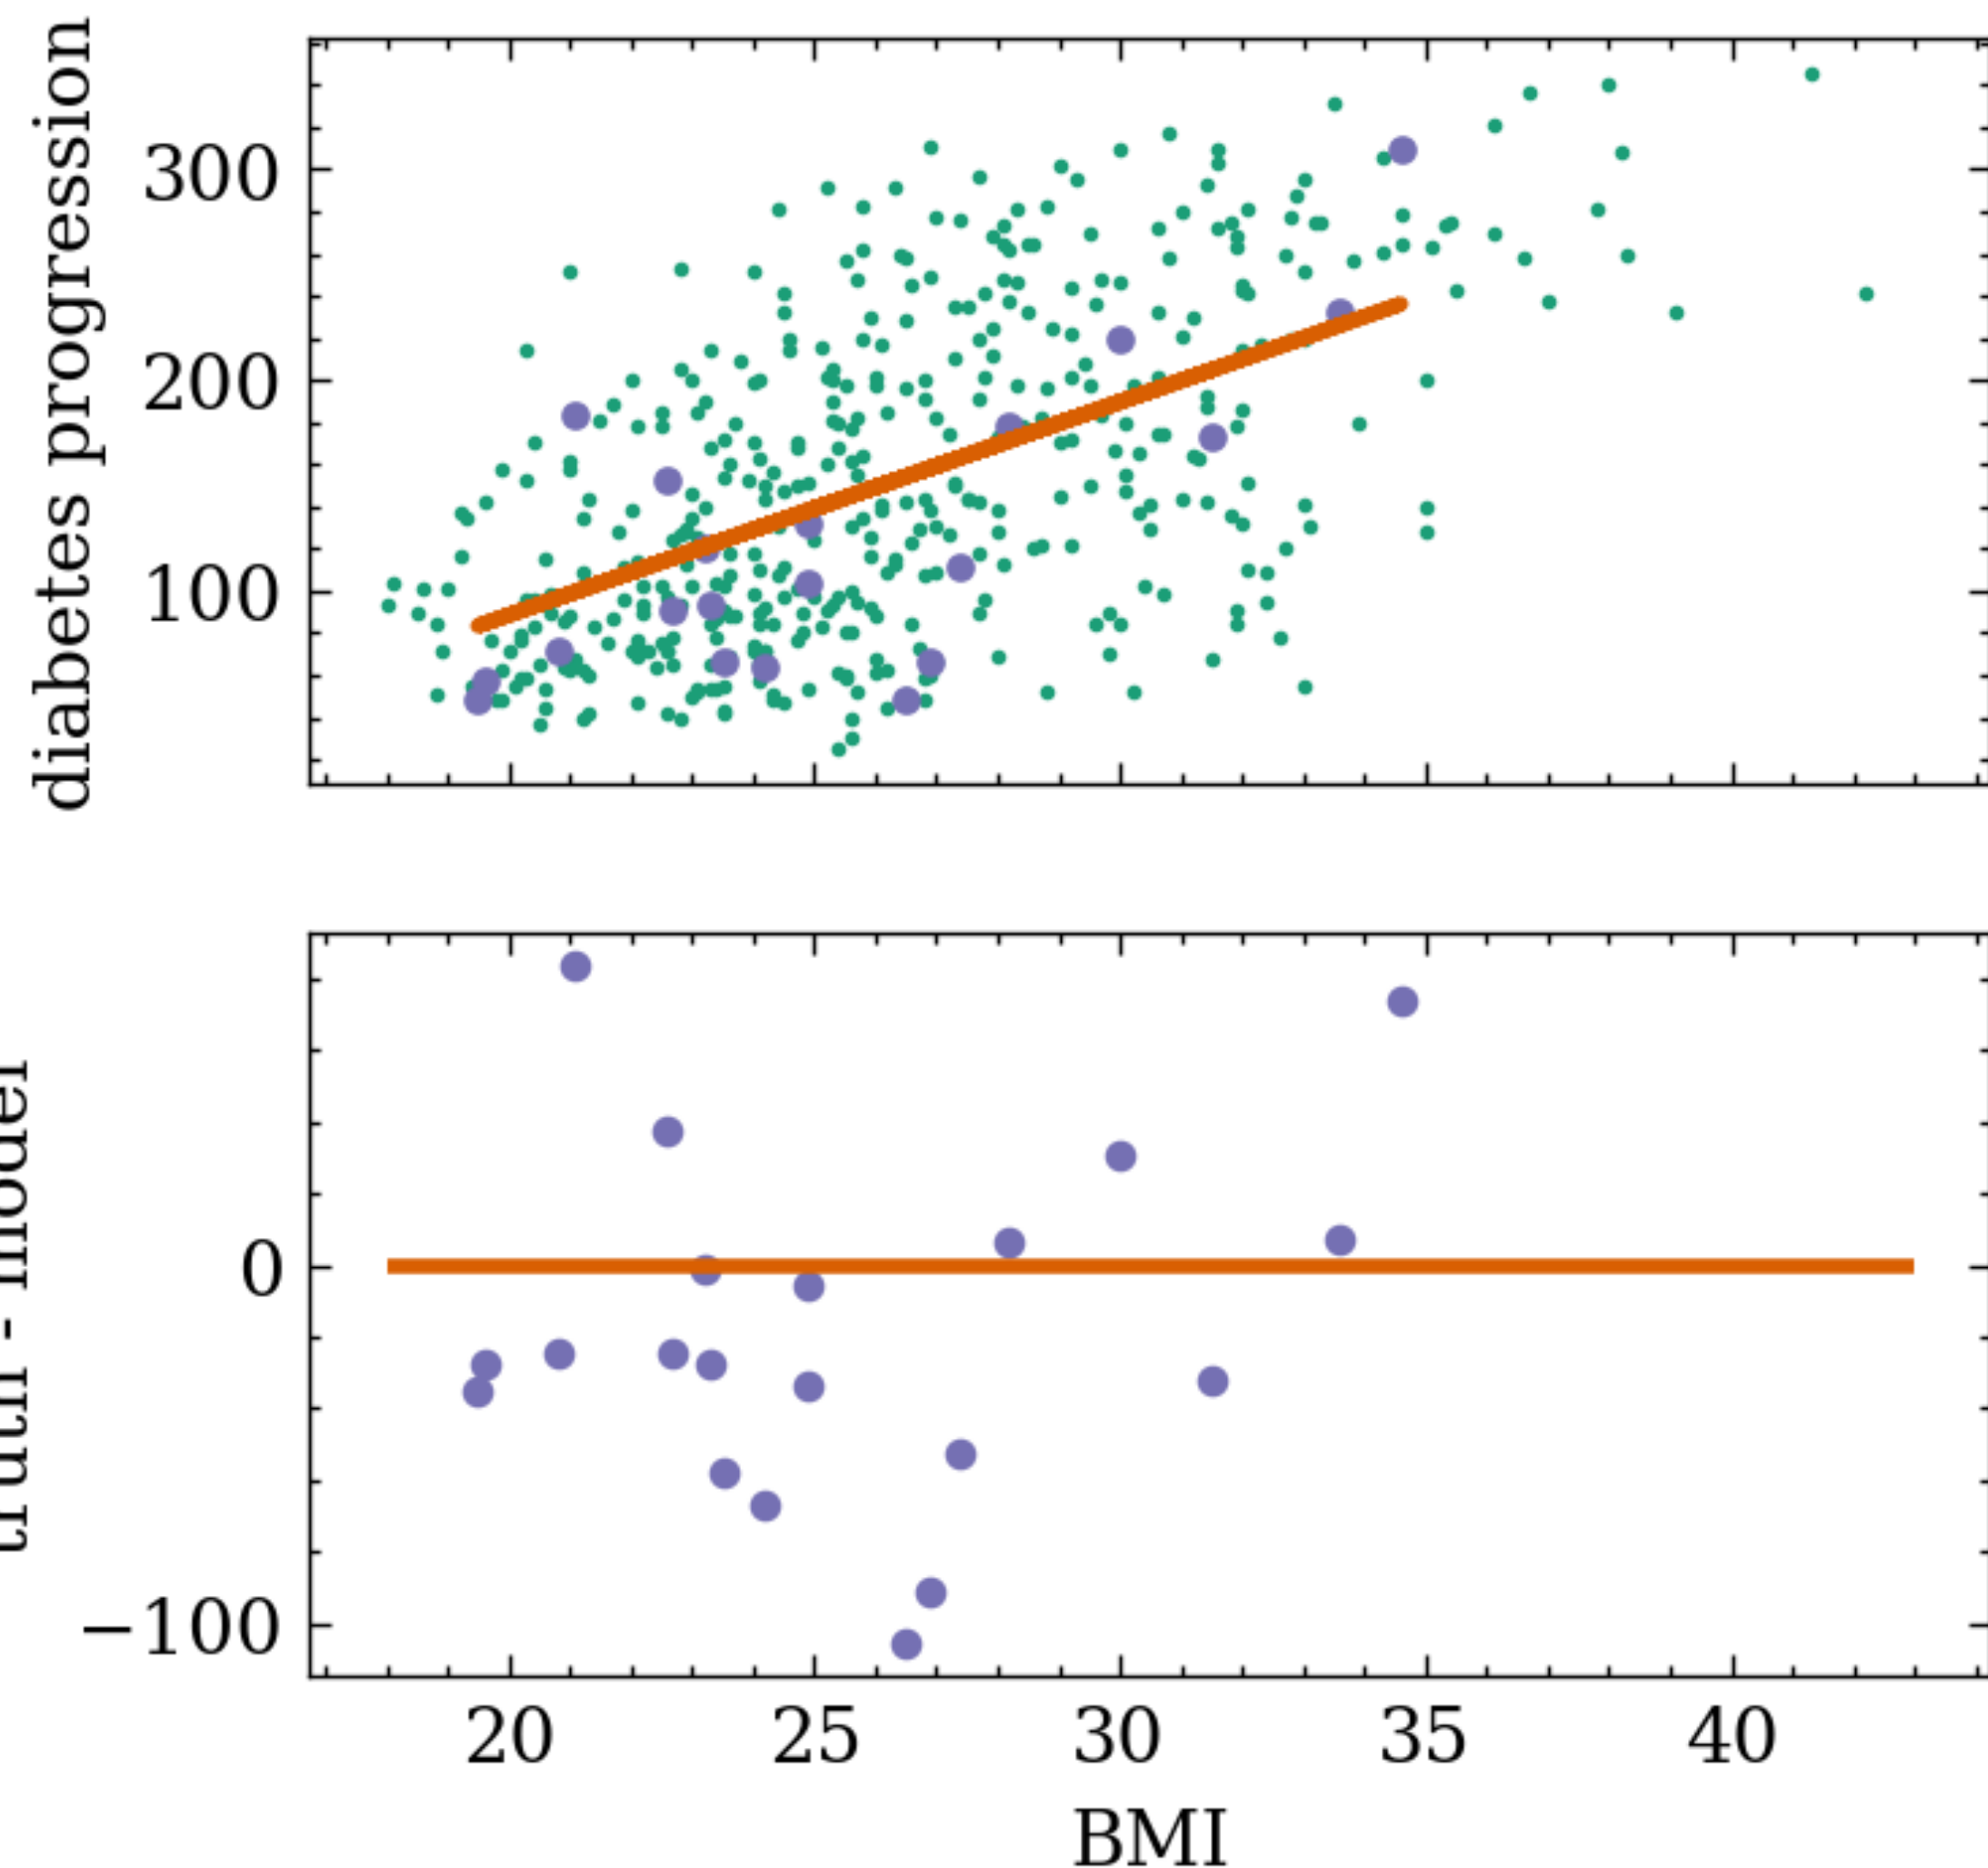
$$y = X\beta$$

## Ordinary Least Squares

Compute the vector  $\beta$  that minimizes:

$$\sum_i^N (y_i - X_i \beta)^2$$

# Linear Regression



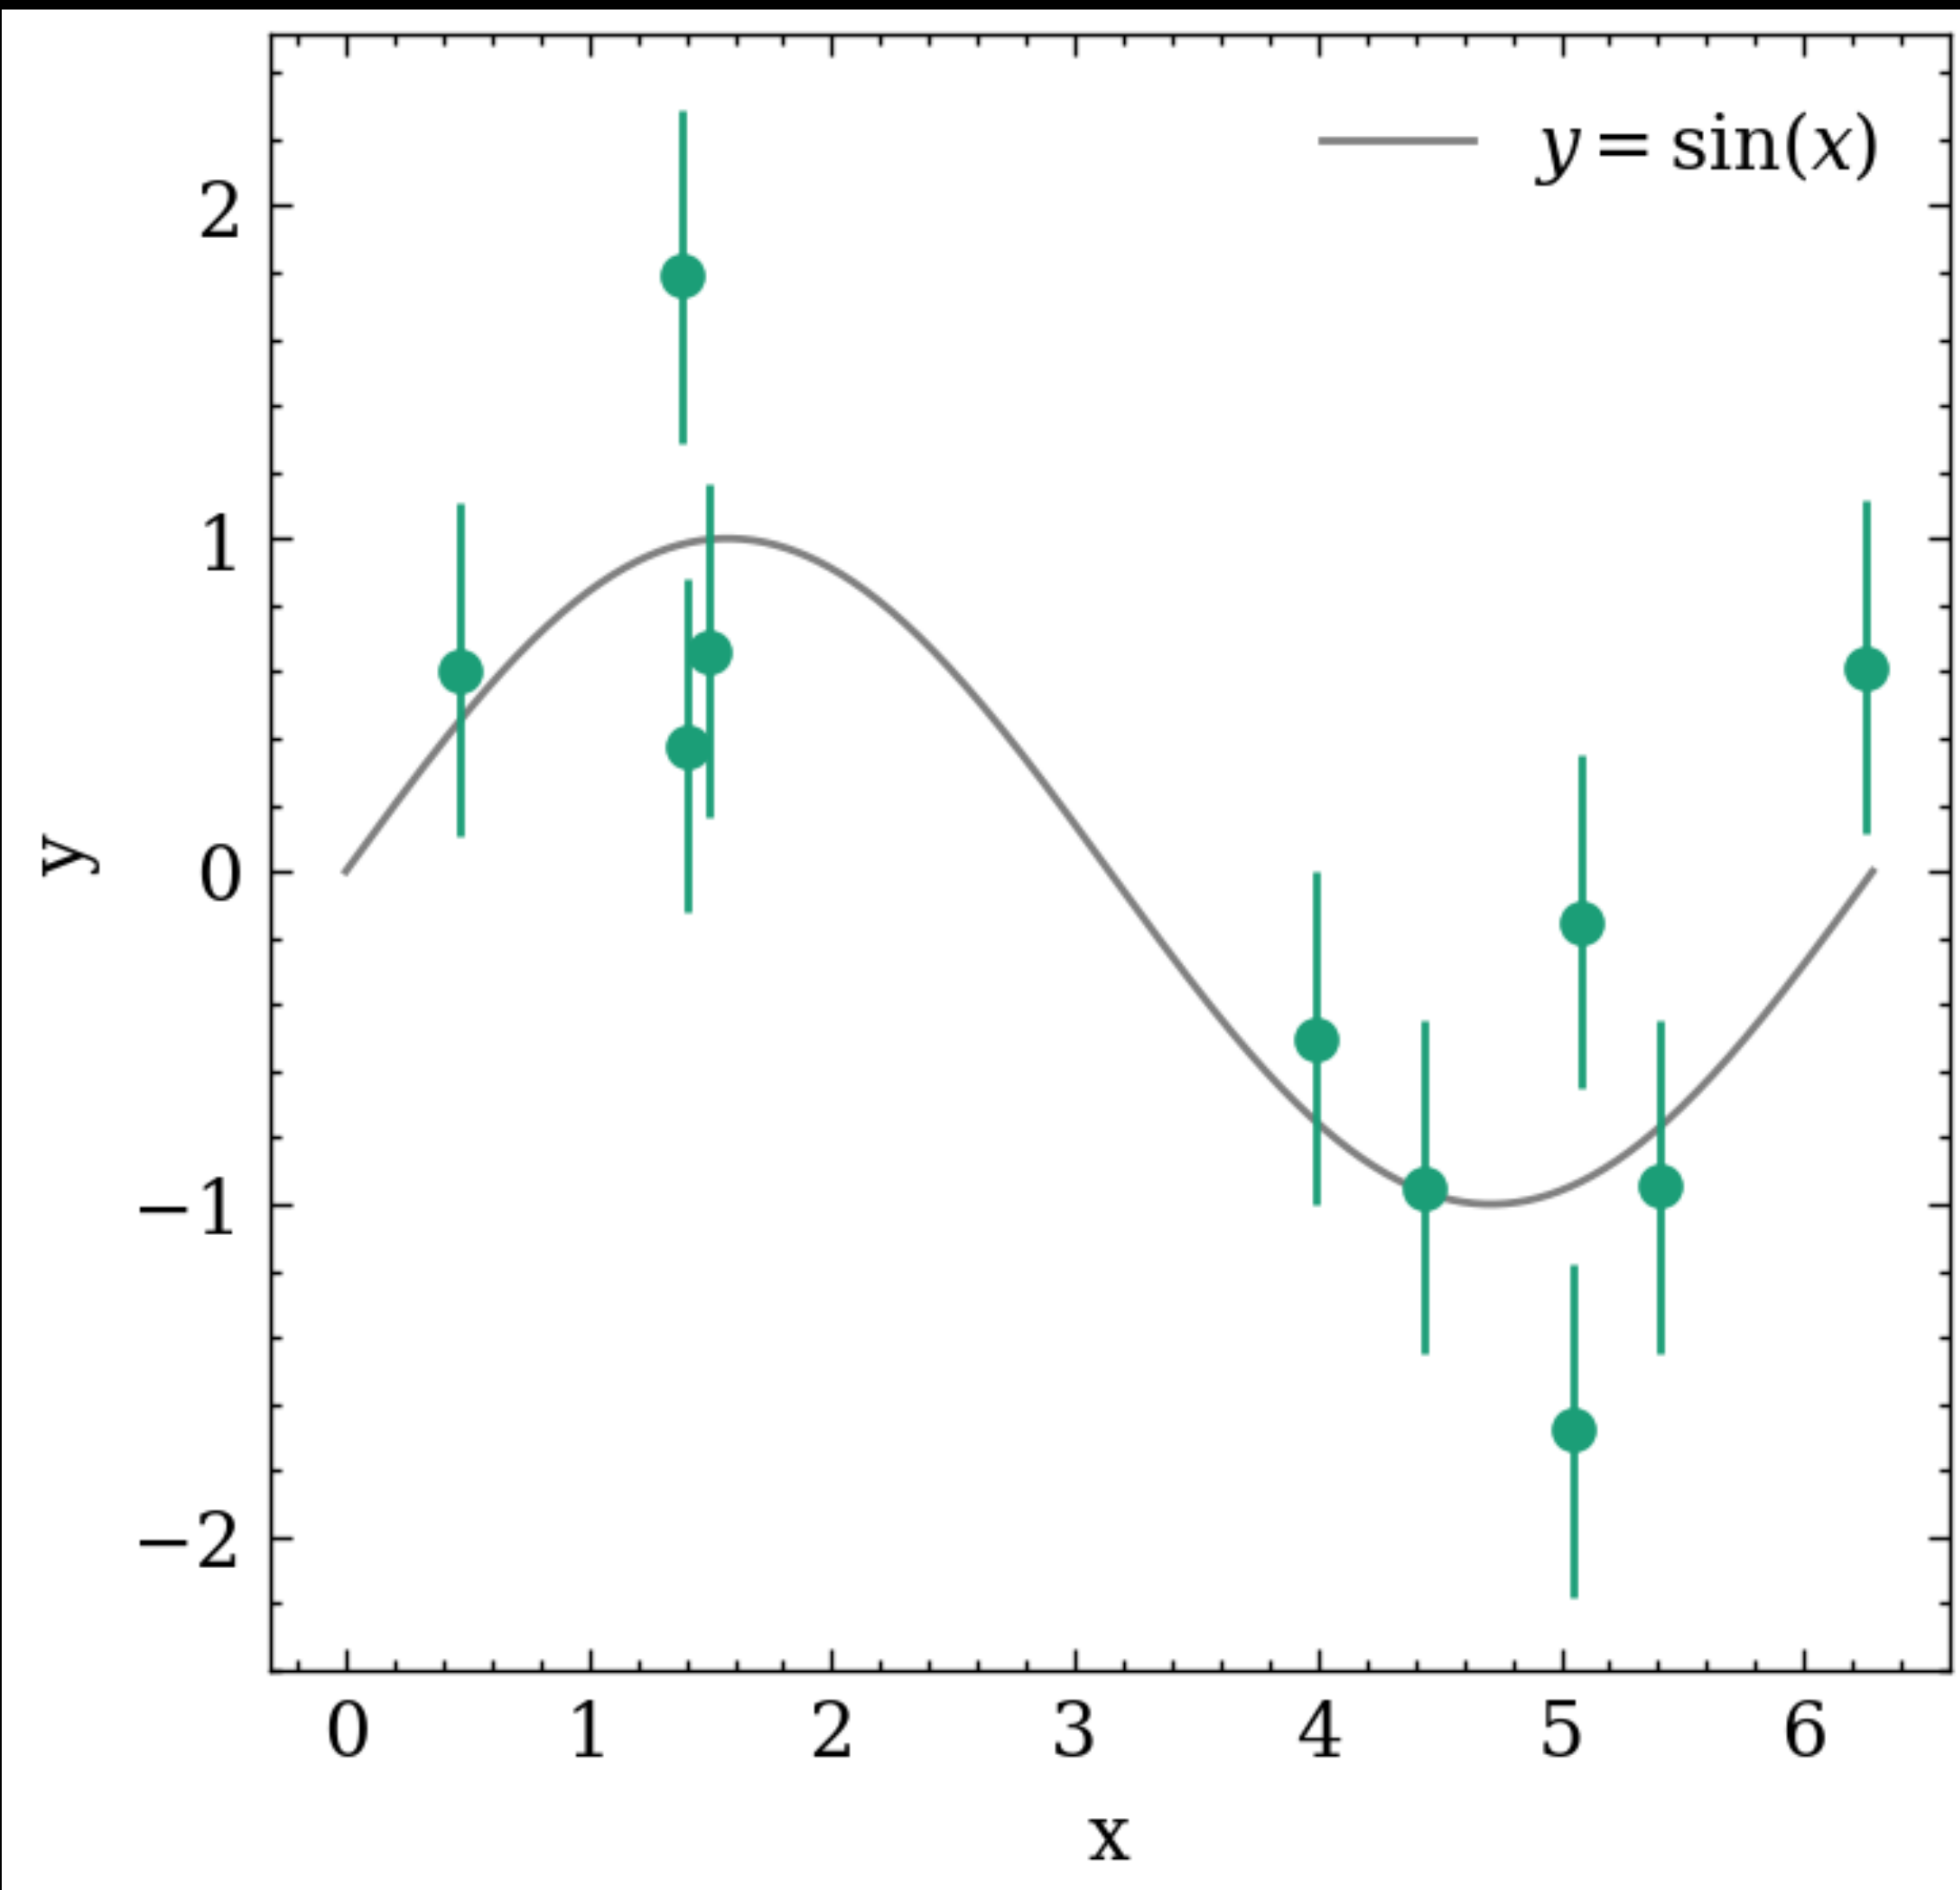
Evaluate the generalizability of the model with a test set of  $M$  input-output pairs:

$$(X_1, y_1), (X_2, y_2), \dots (X_M, y_M)$$

Test set data is not included in the training set.



# Gaussian Process Regression



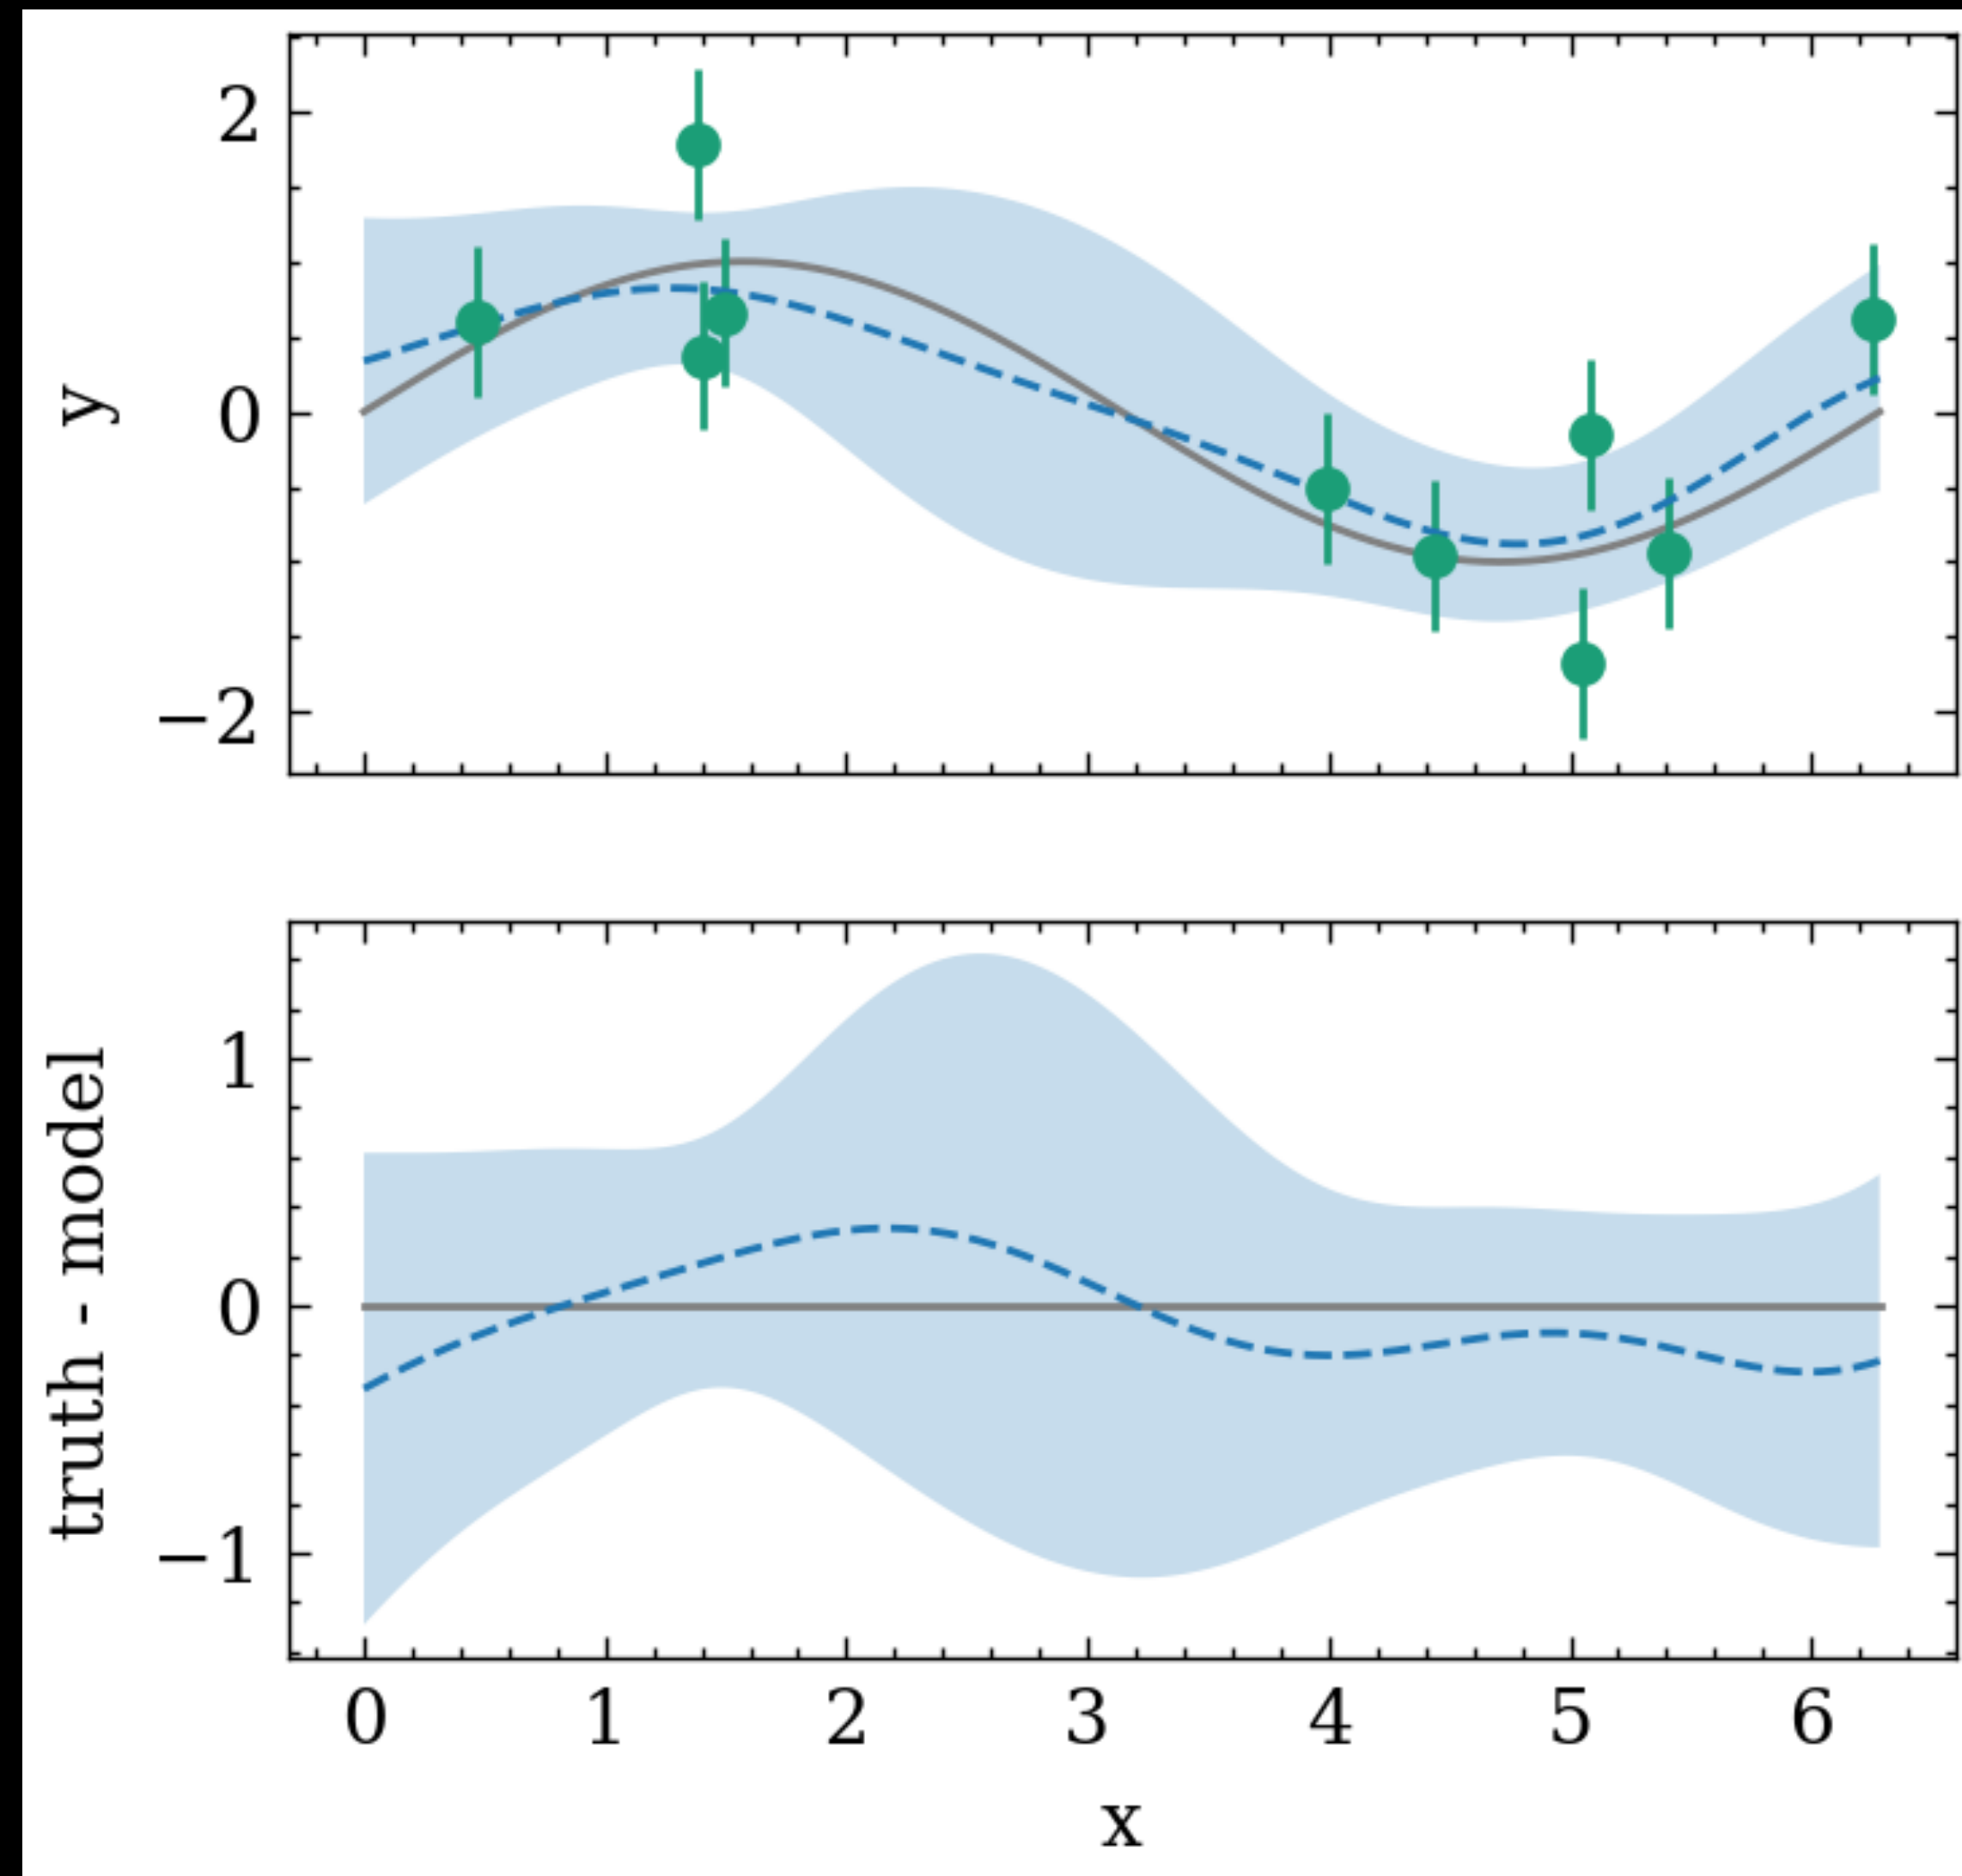
input-output pairs:

- $x_i$
- $y_i = \sin(x_i) + \text{Gaussian noise}$

Goal: learn a function  $h$  that approximates the true function  $f$ :

$$\hat{y} = h(x)$$

# Gaussian Process Regression



A regression method where the prediction is probabilistic

- Gaussian
- compute empirical confidence intervals

`sklearn.gaussian_process.GaussianProcessRegressor`

Carl Eduard Rasmussen and Christopher K.I. Williams,  
“Gaussian Processes for Machine Learning”, MIT Press  
2006.

Examples of problems in your area of research  
where *regression* could be used?

# Common algorithms

Supervised	Unsupervised
Regression	Clustering
Classification	Dimensionality Reduction

# Classification

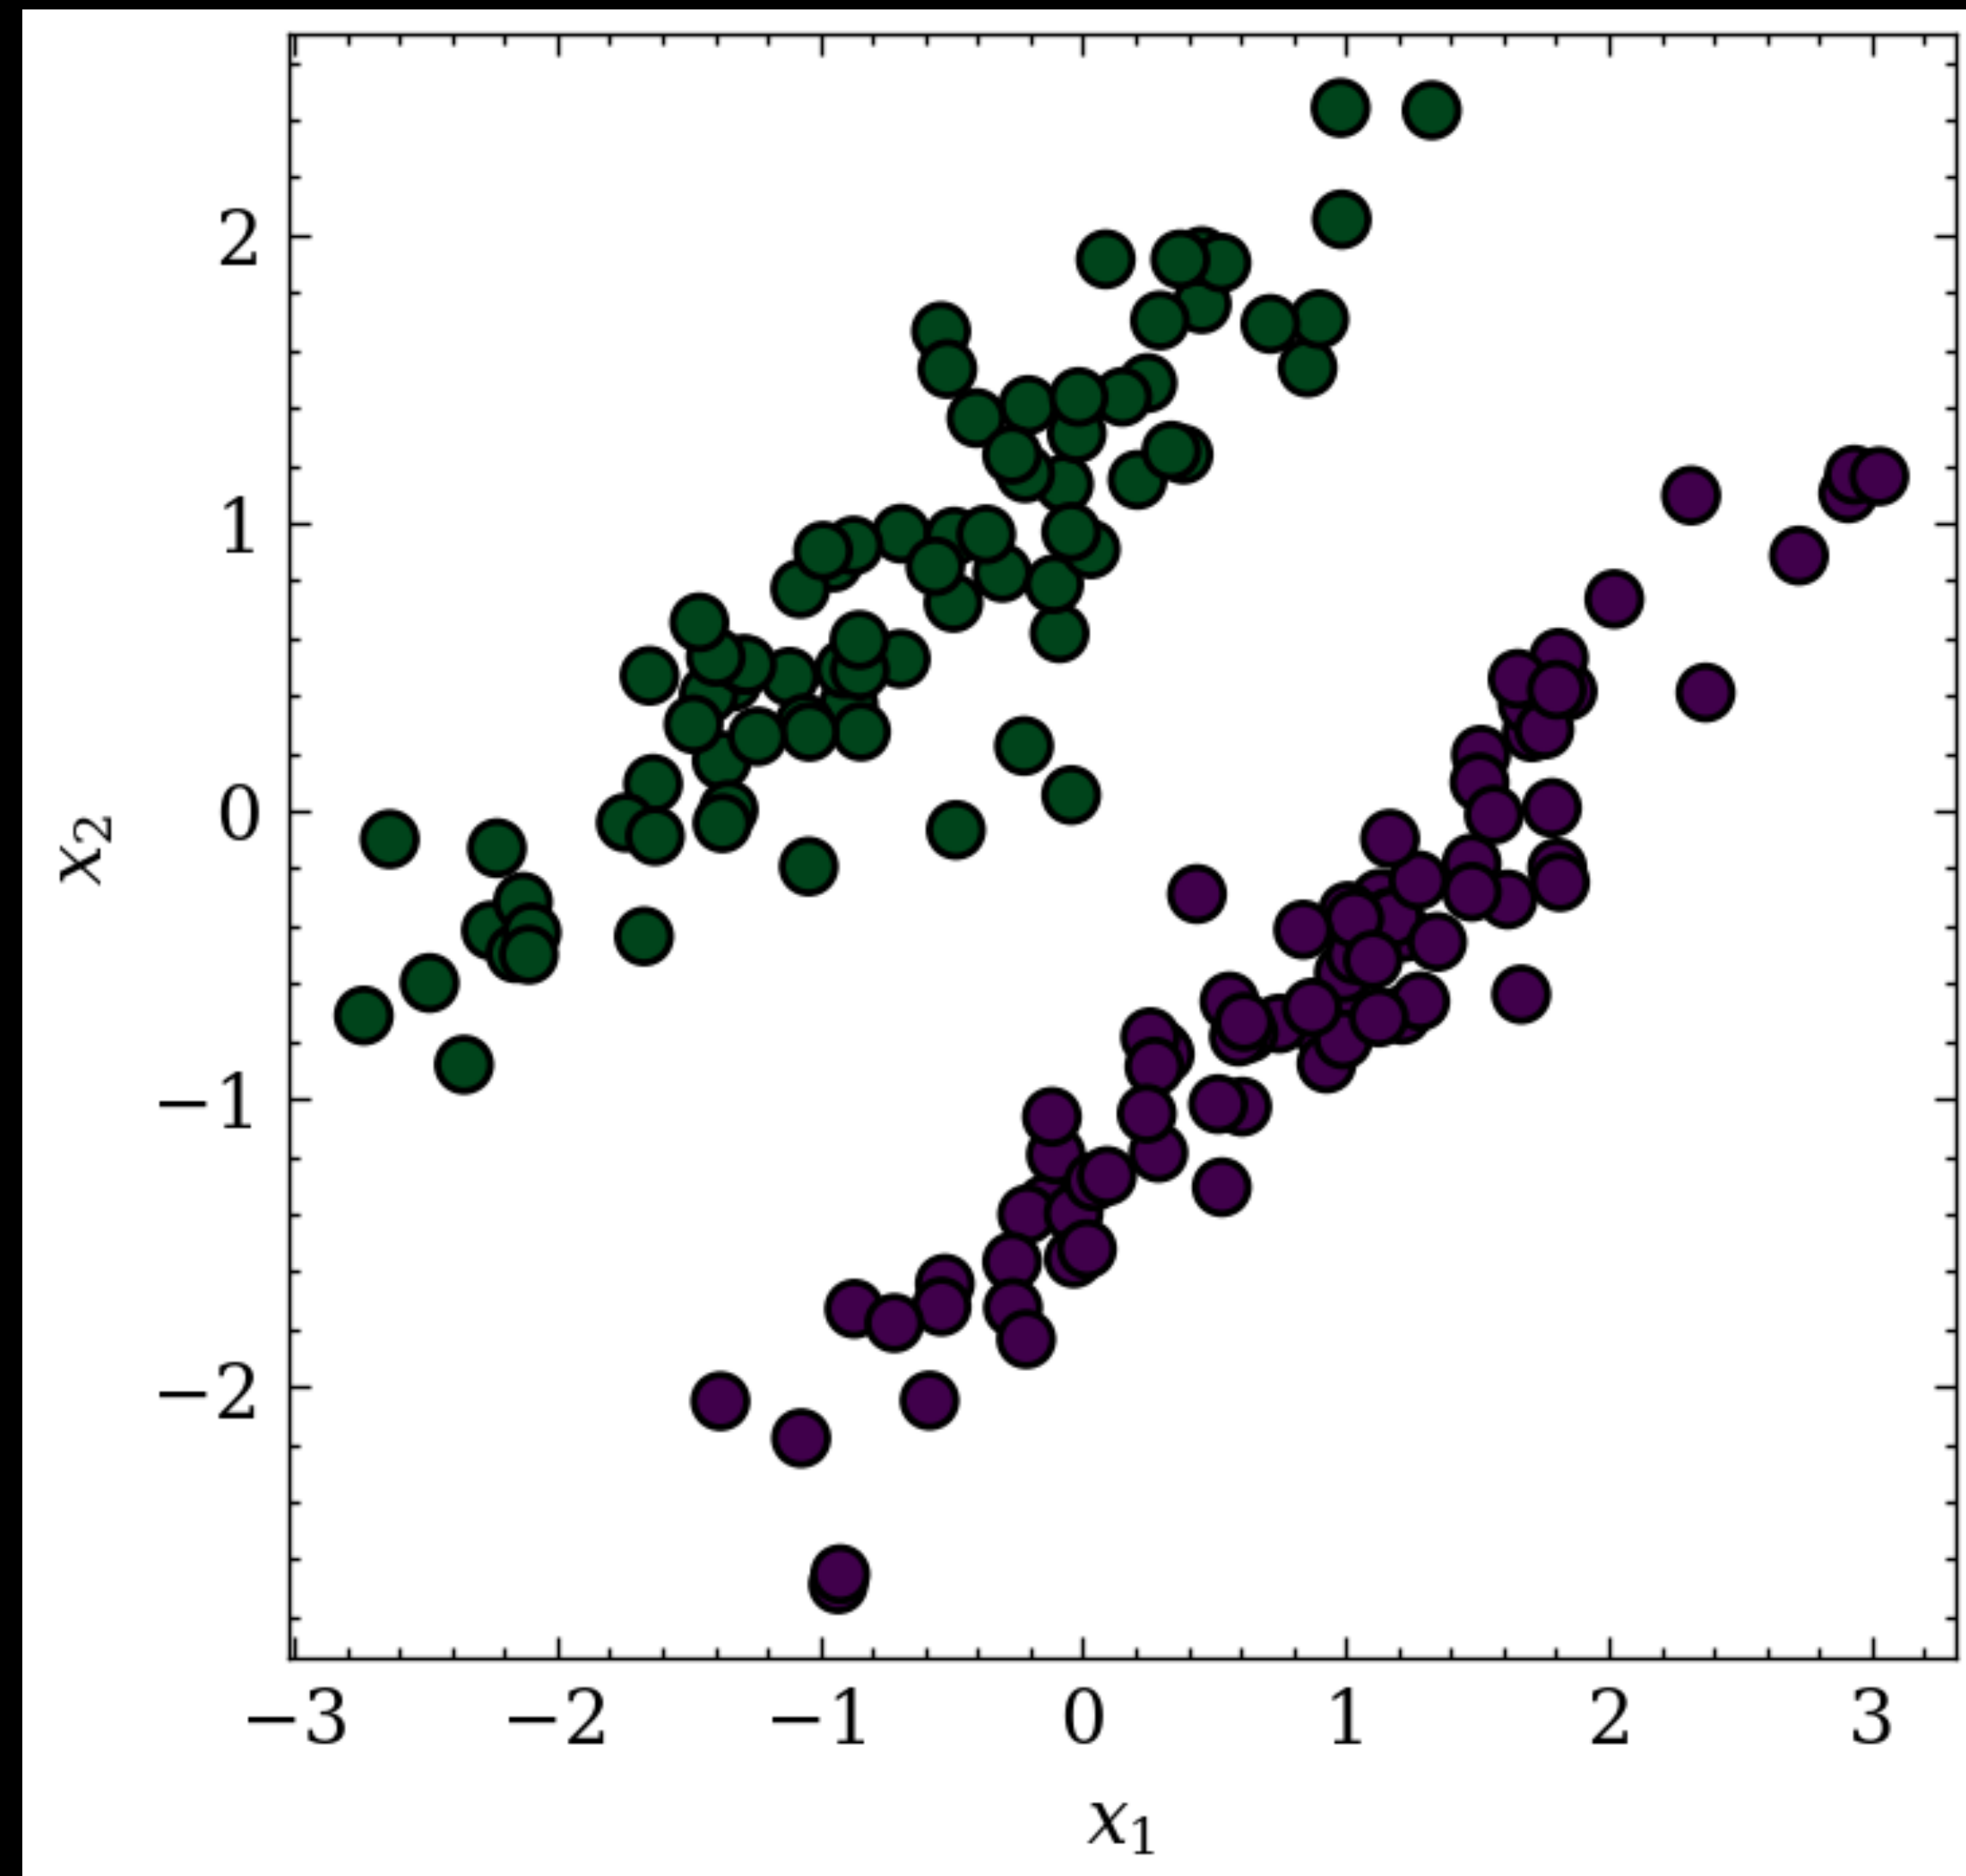
Training set of  $N$  input-output pairs:

$$(X_1, y_1), (X_2, y_2), \dots (X_N, y_N)$$

where desired output value is a discrete value.

Goal: learn a function  $h$  that approximates the true function  $f$ .

# Classification



input-output pairs:

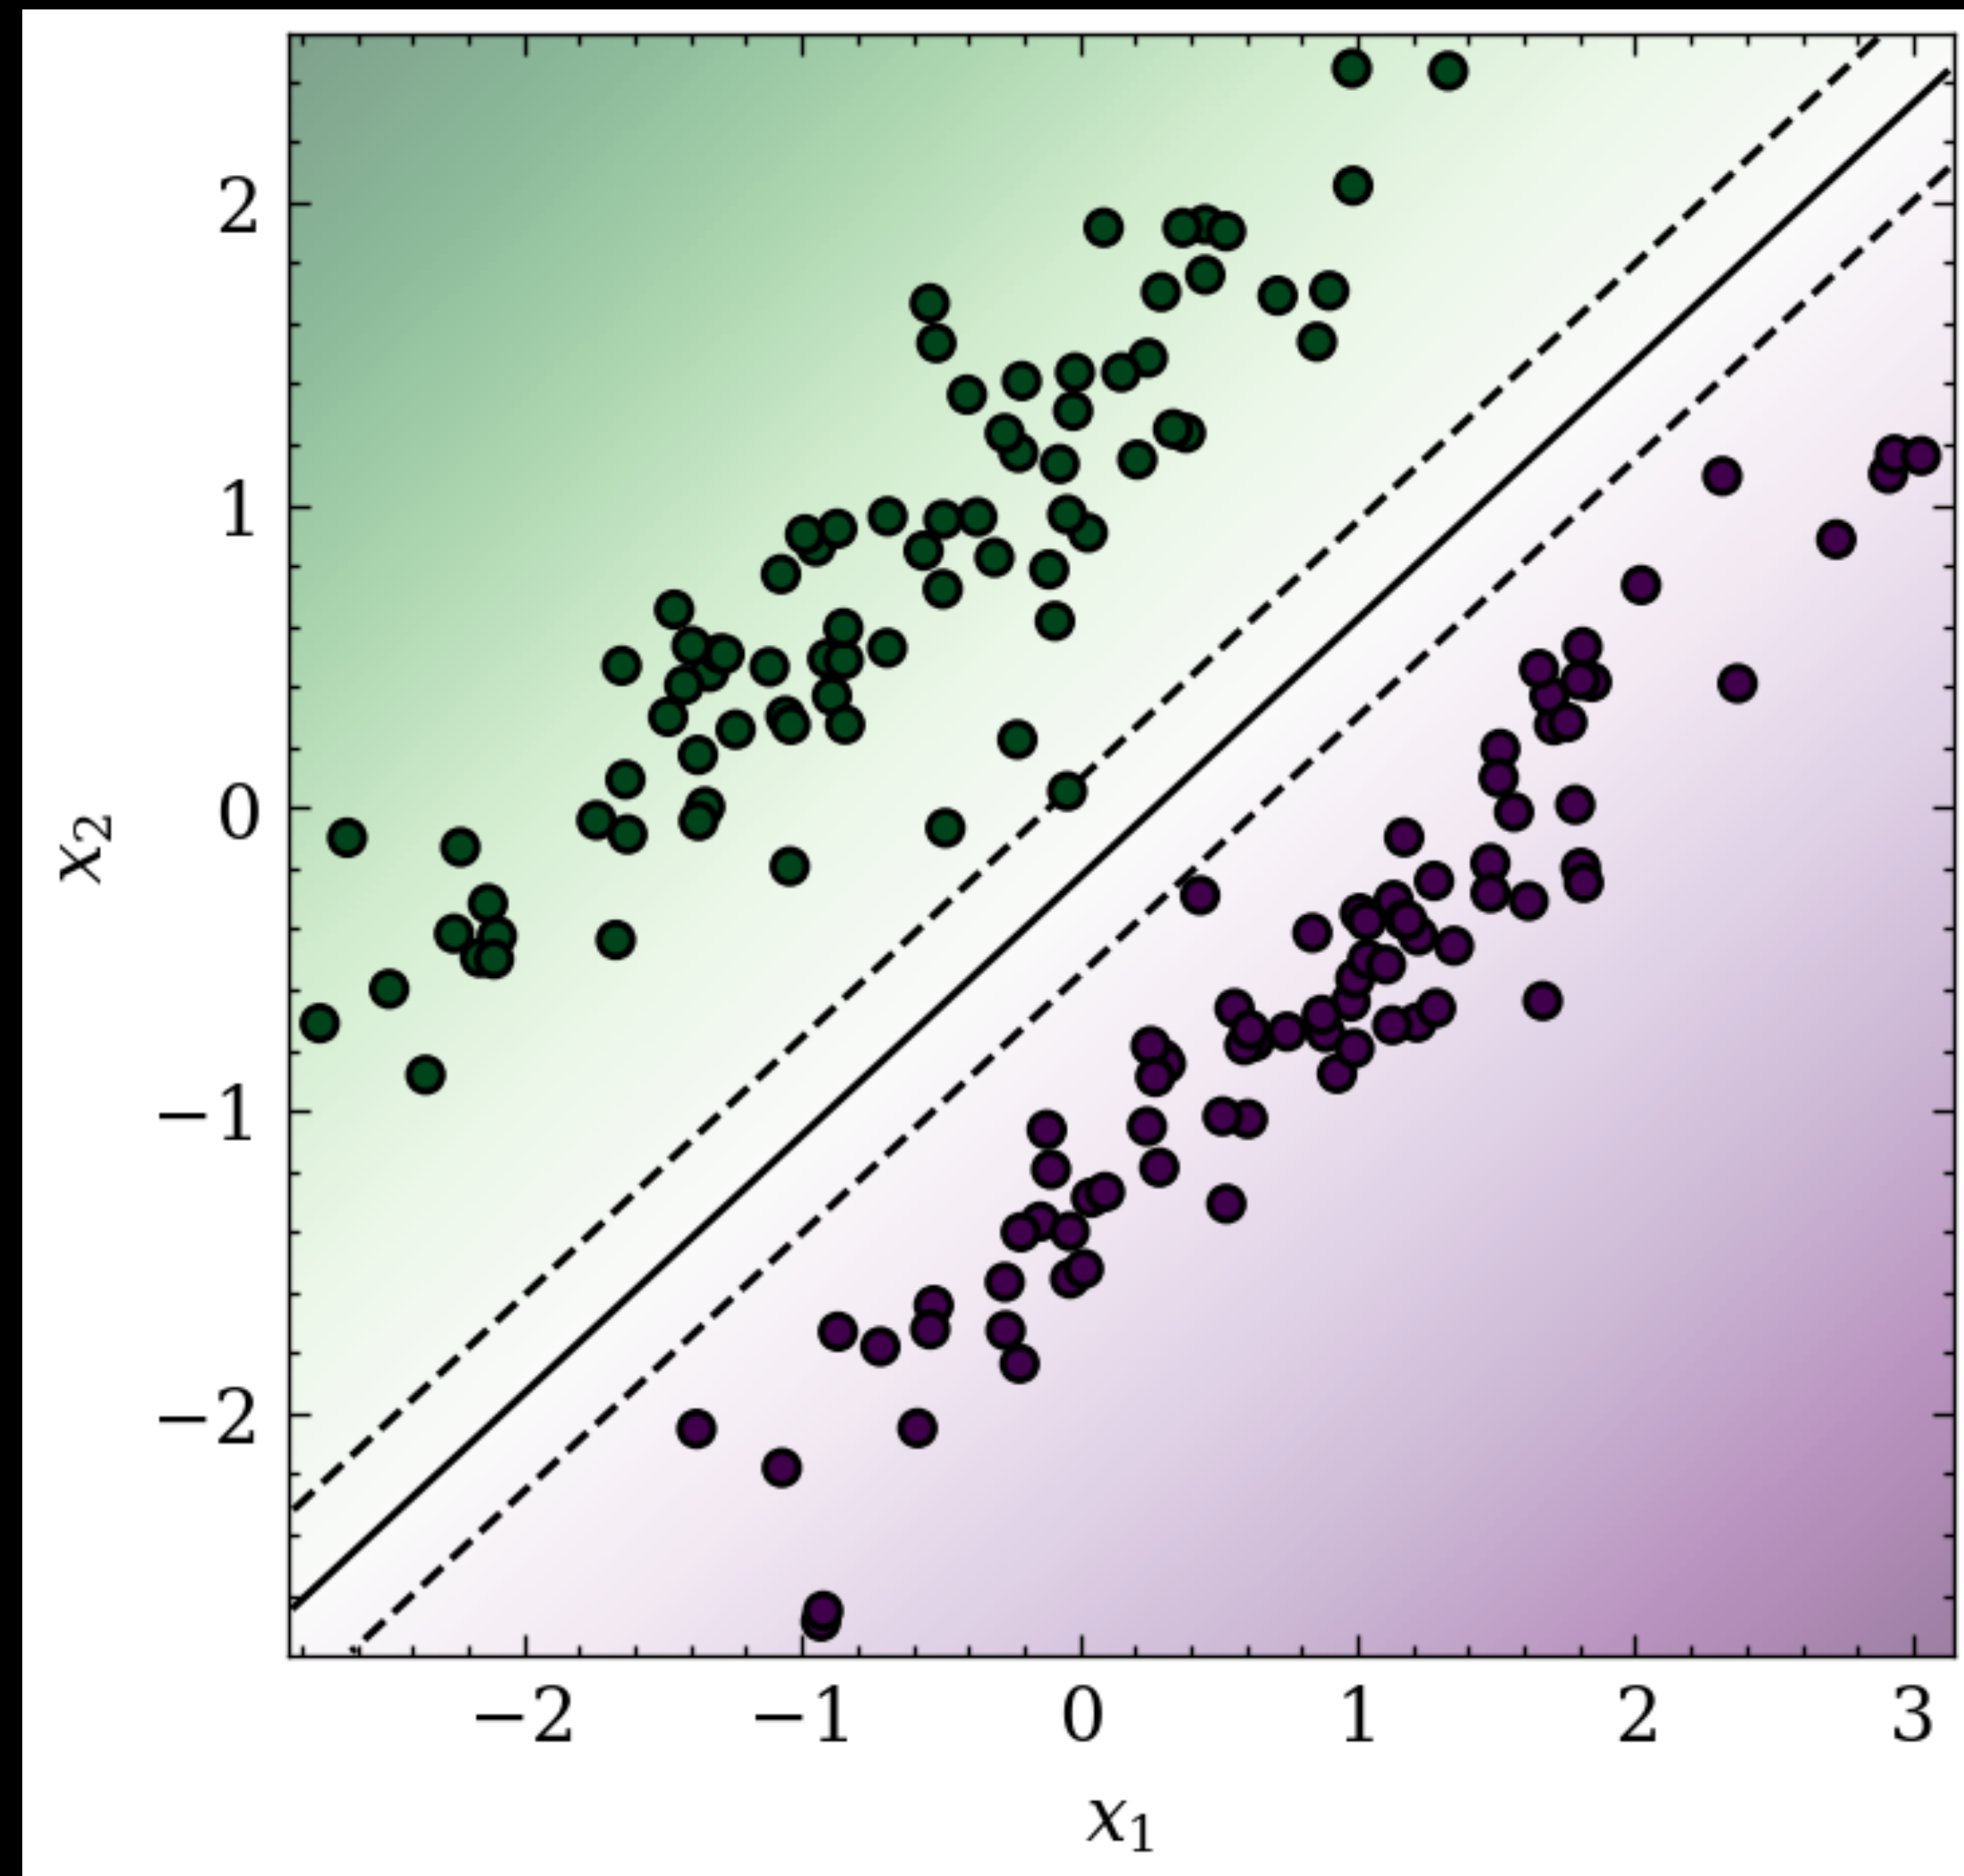
- $X_i = \{x_{i,1}, x_{i,2}\}$
- $y_i$ , class label, either 1 or -1

Goal: learn a function  $h$  that approximates the true function  $f$ :

*predicted class label,  $\hat{y} = h(x_1, x_2)$*



# Support Vector Classification



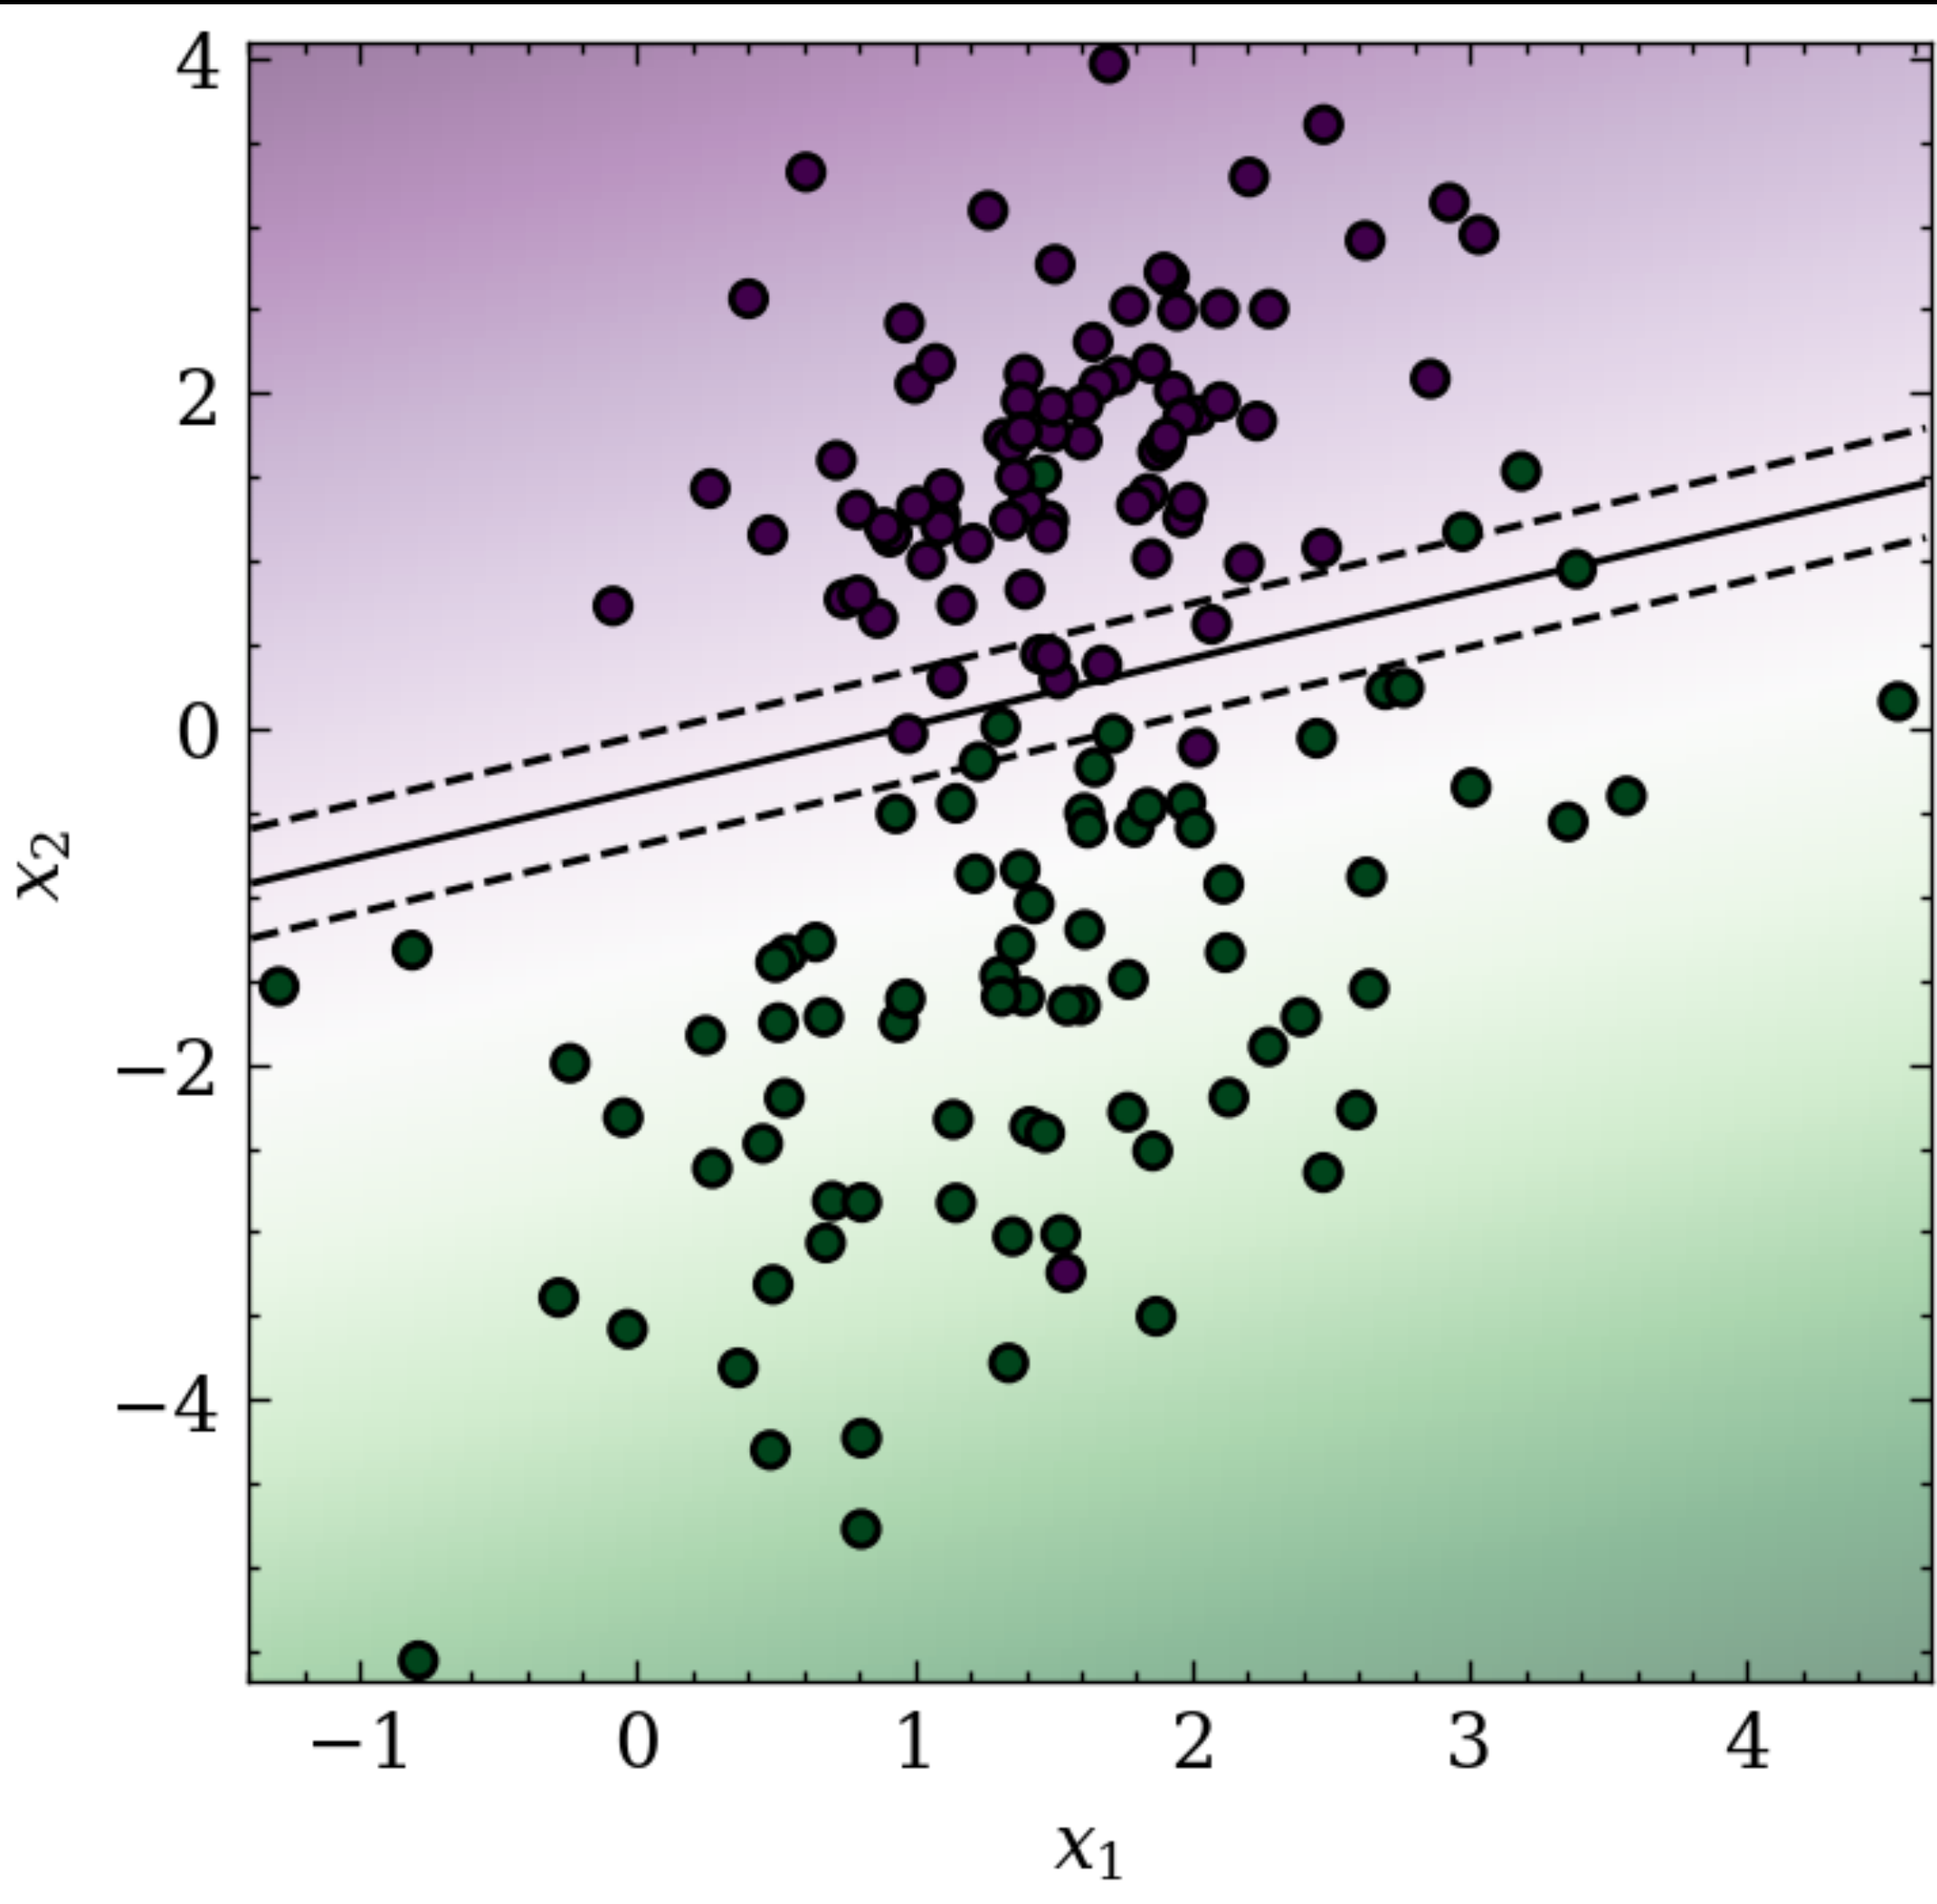
Calculate the line that separates the classes with the maximal margin.

Margin: area between two parallel lines that separate the two classes of data.

- 3-D : plane
- > 3-D : hyperplane

*sklearn.svm.SVC*

# Support Vector Classification



What if the data are not linearly separable?

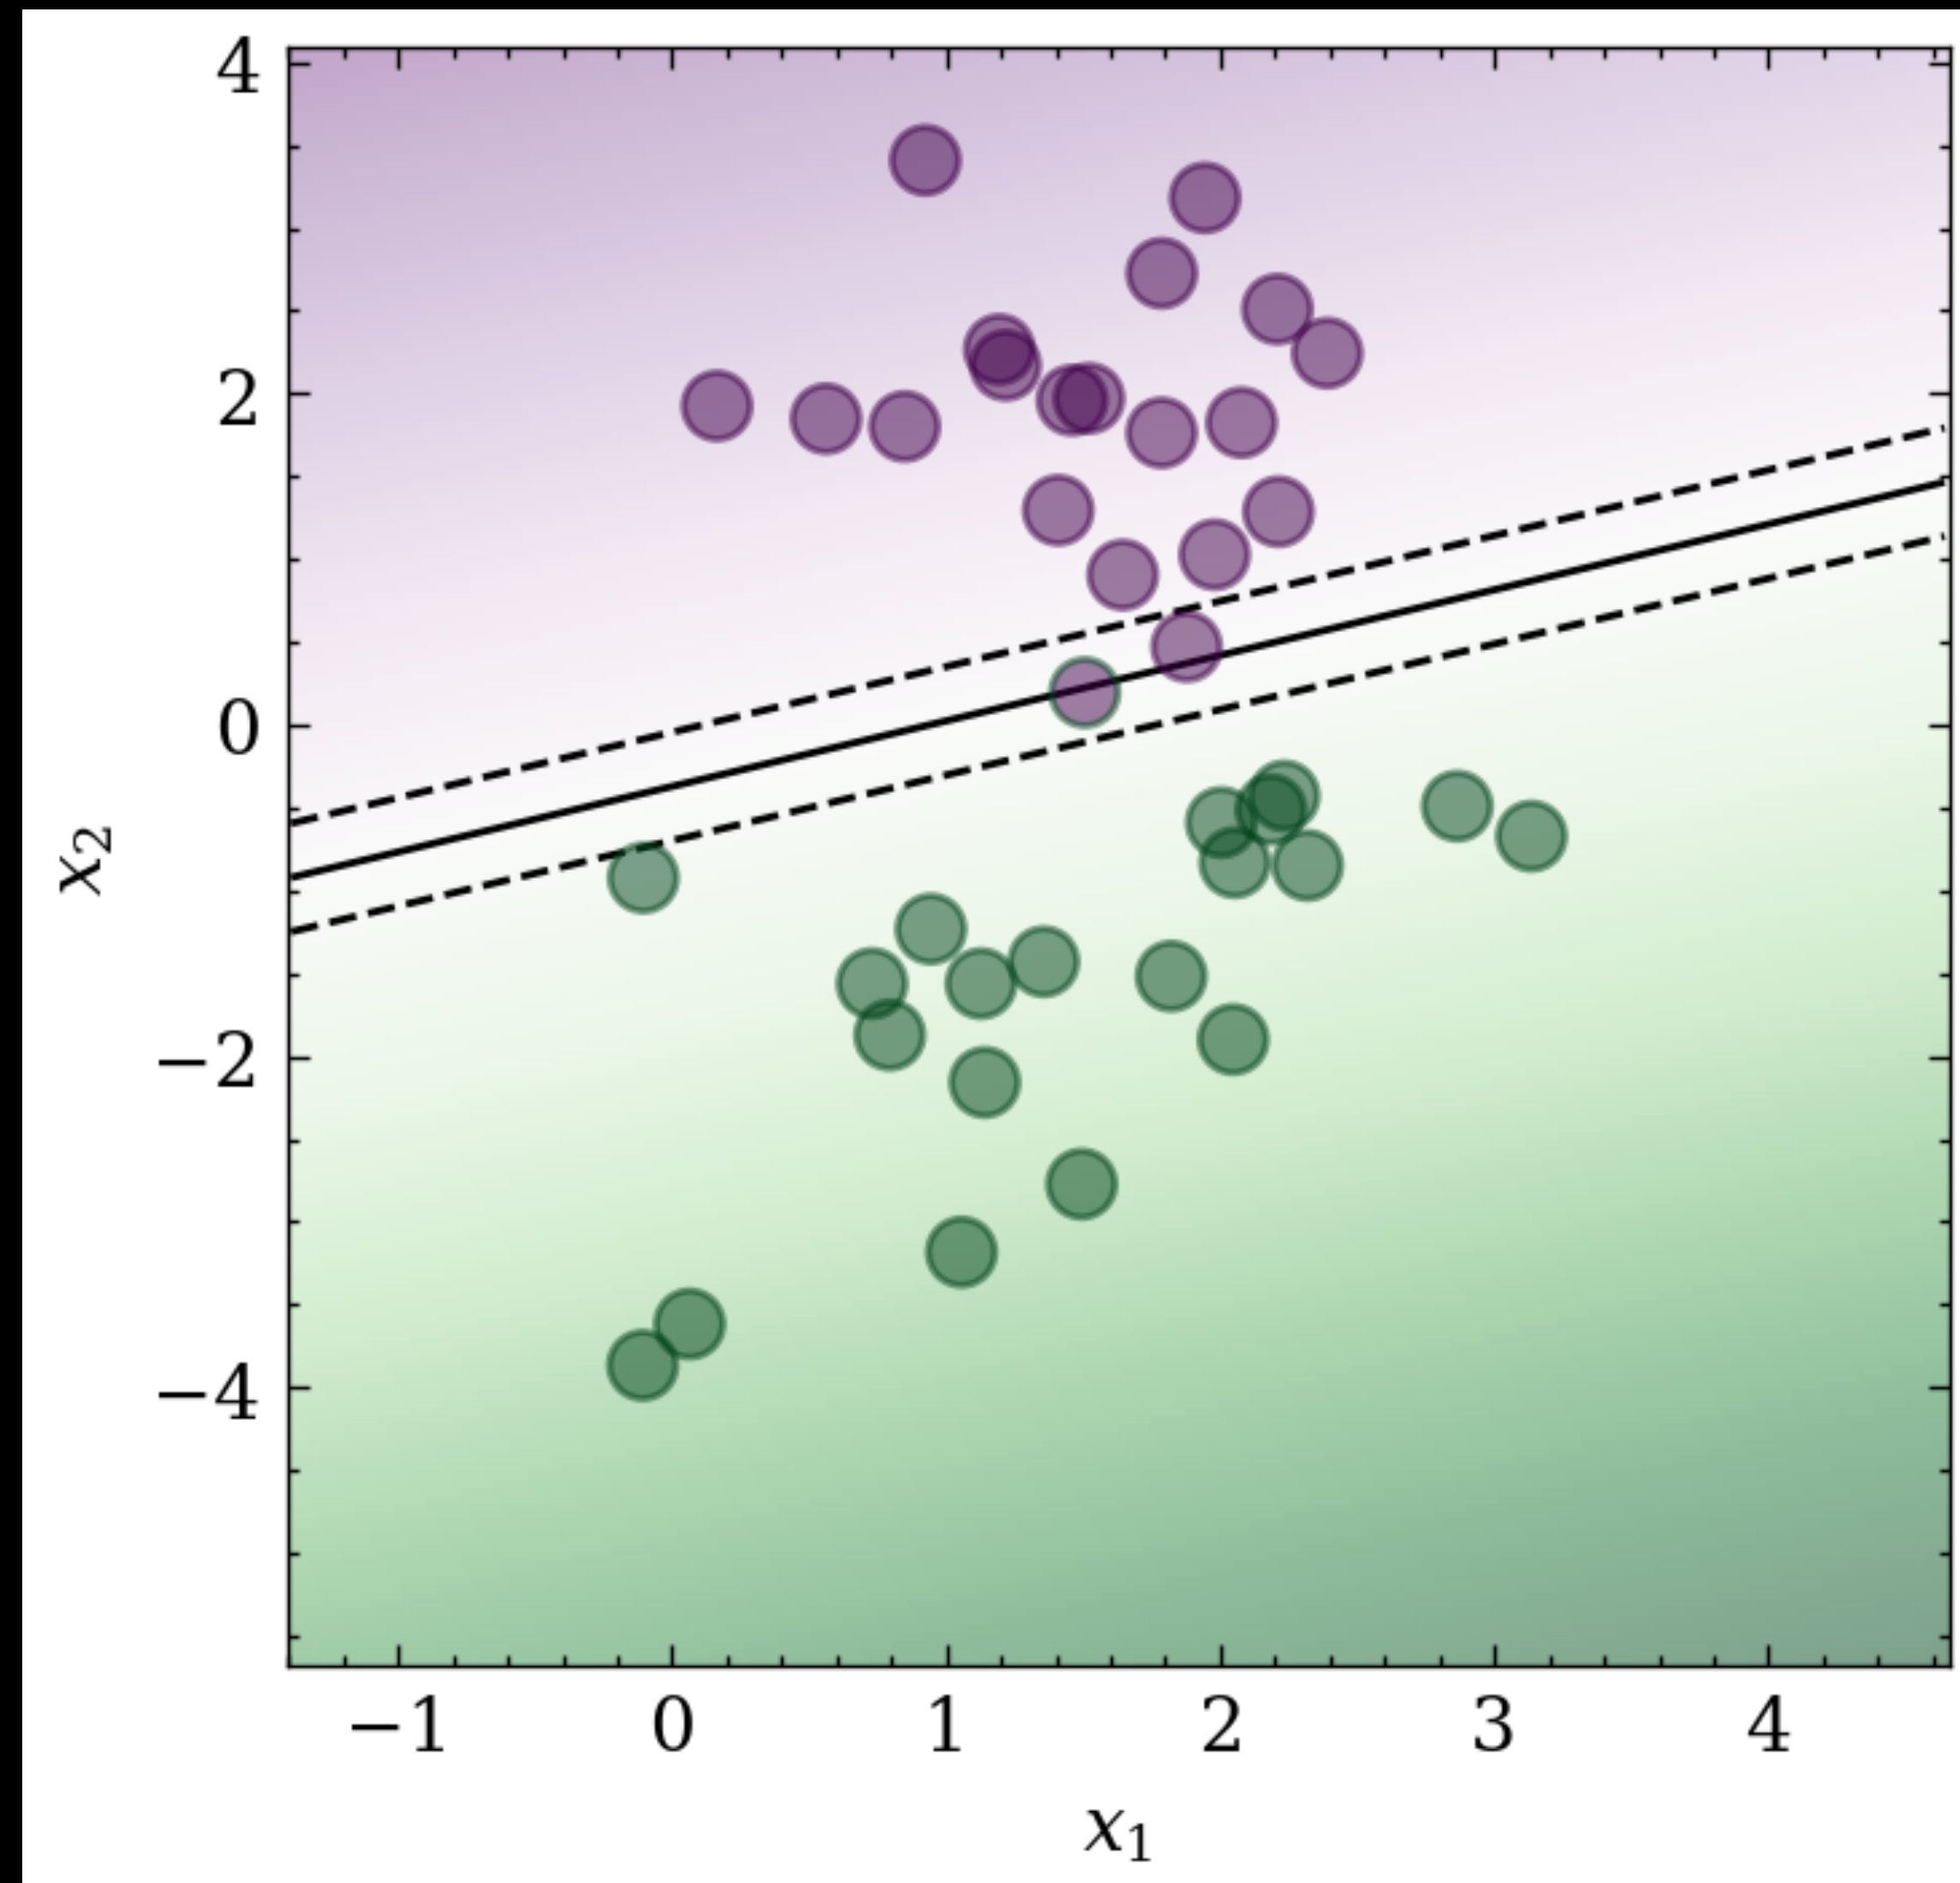
Minimize a function that is a sum over values for all training data:

- 0 - if on correct side of margin
- value proportional to the distance from the margin

*sklearn.svm.SVC*



# Support Vector Classification



Evaluate the generalizability of the model with a test set of  $M$  input-output pairs:

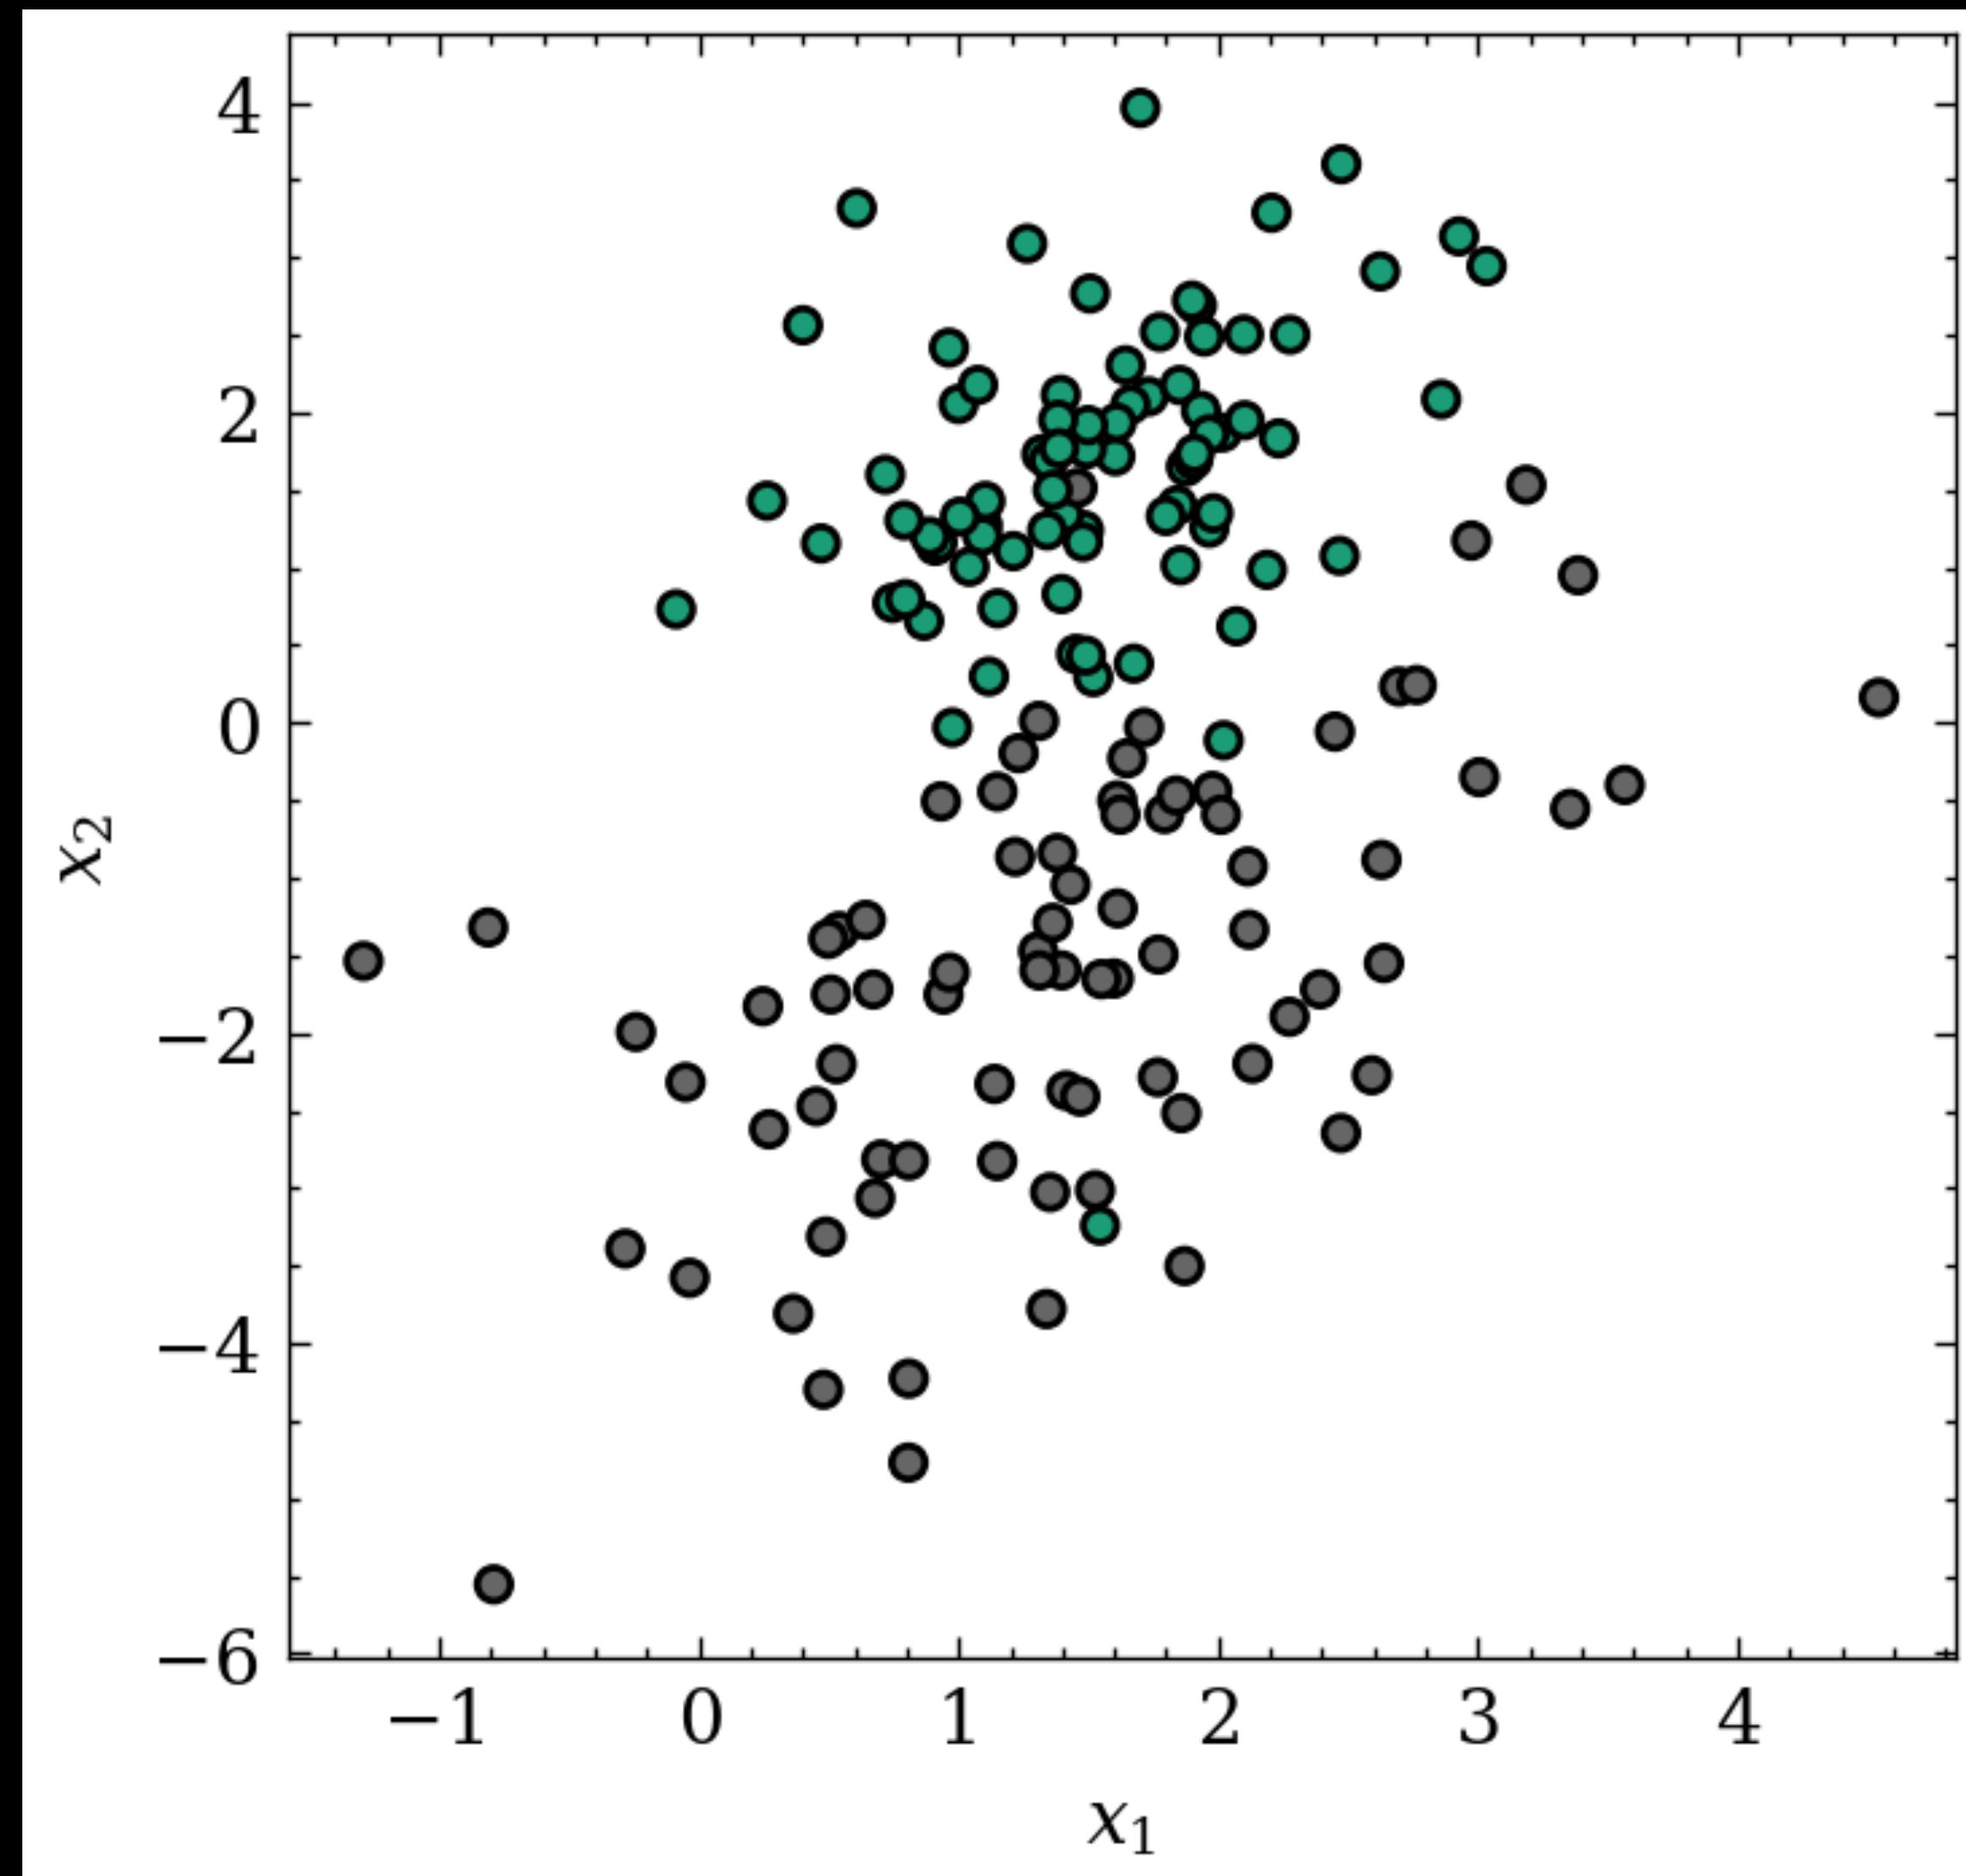
$$(X_1, y_1), (X_2, y_2), \dots (X_M, y_M)$$

Test set data is not included in the training set.

$$\text{Accuracy} = \frac{1}{M} \sum_i^M (\hat{y}_i = y_i)$$

`sklearn.svm.SVC`

# K-Nearest Neighbors

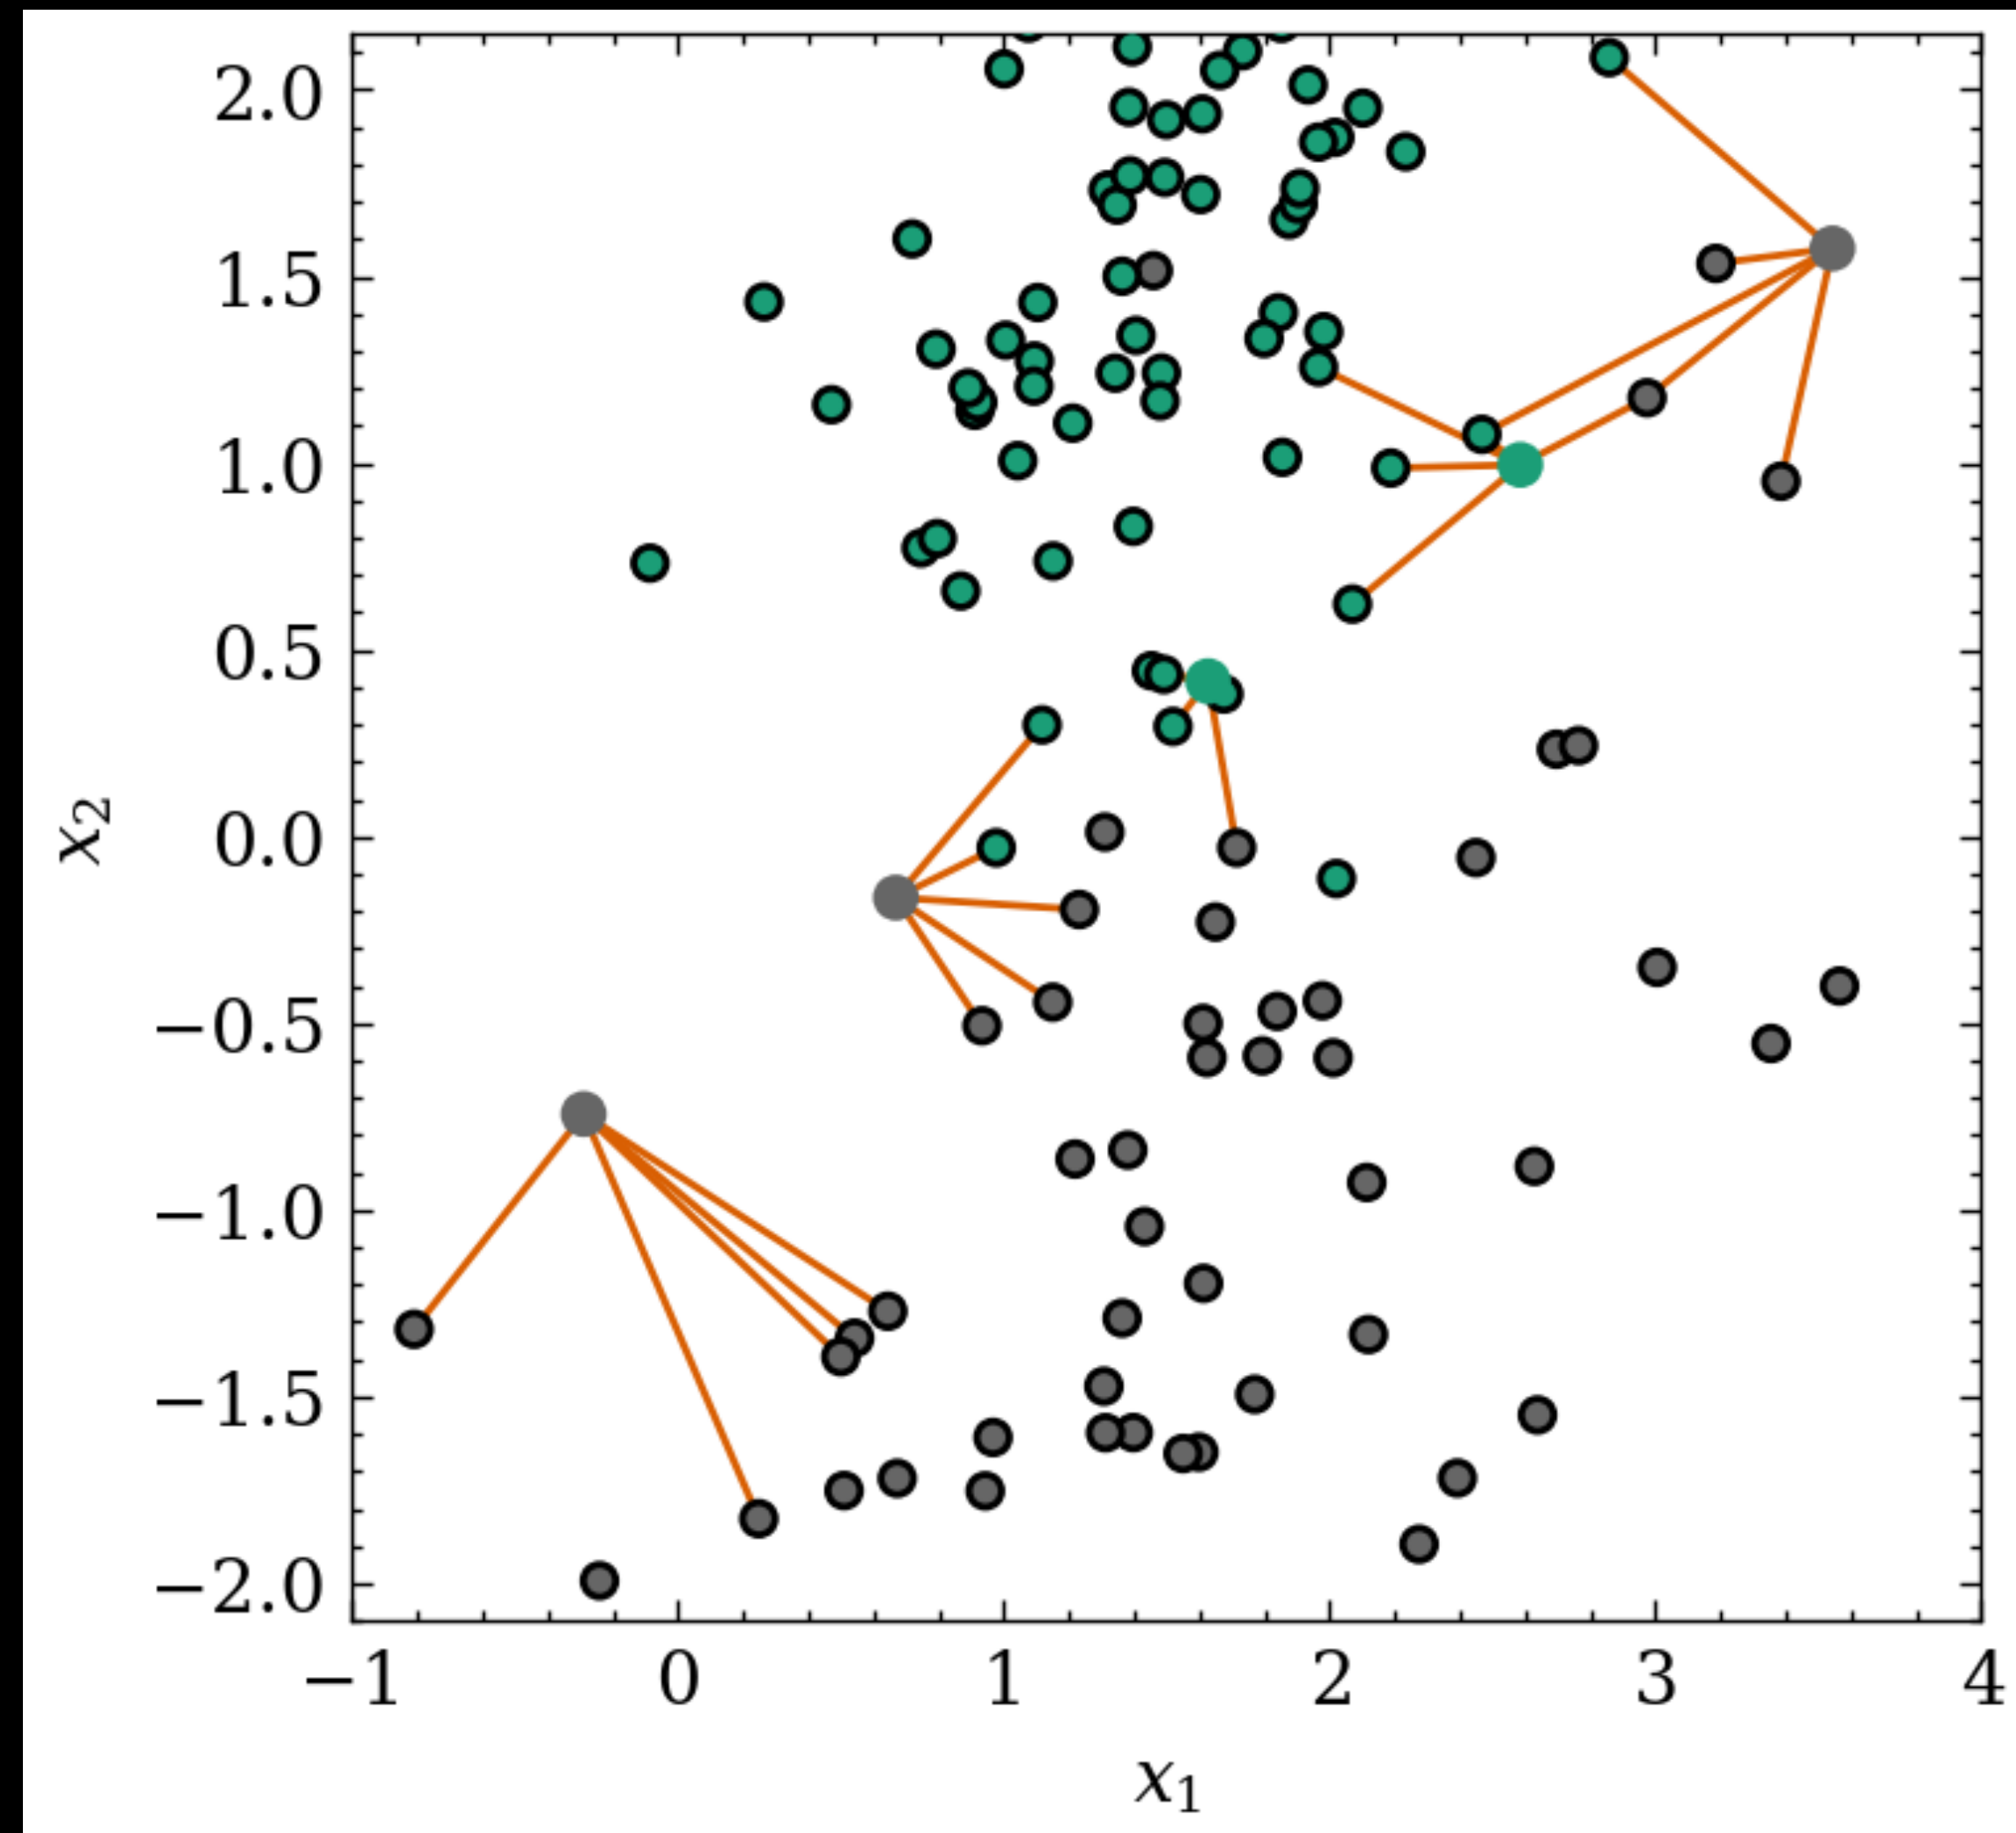


Classification is computed from a simple majority vote of the nearest neighbors of each point.

Does not construct a general model; instead stores the training data.

*`sklearn.neighbors.KNeighborsClassifier`*

# K-Nearest Neighbors



Classification is computed from a simple majority vote of the nearest neighbors of each point.

Does not construct a general model; instead stores the training data.

At left, the case where  $k = 5$ .

*`sklearn.neighbors.KNeighborsClassifier`*

Examples of problems in your area of research  
where *classification* could be used?

# Common algorithms

Supervised	Unsupervised
Regression	Clustering
Classification	Dimensionality Reduction

# Clustering

Training set of N inputs:

$X_1, X_2, \dots, X_N.$

Goal: identify M clusters in the data.



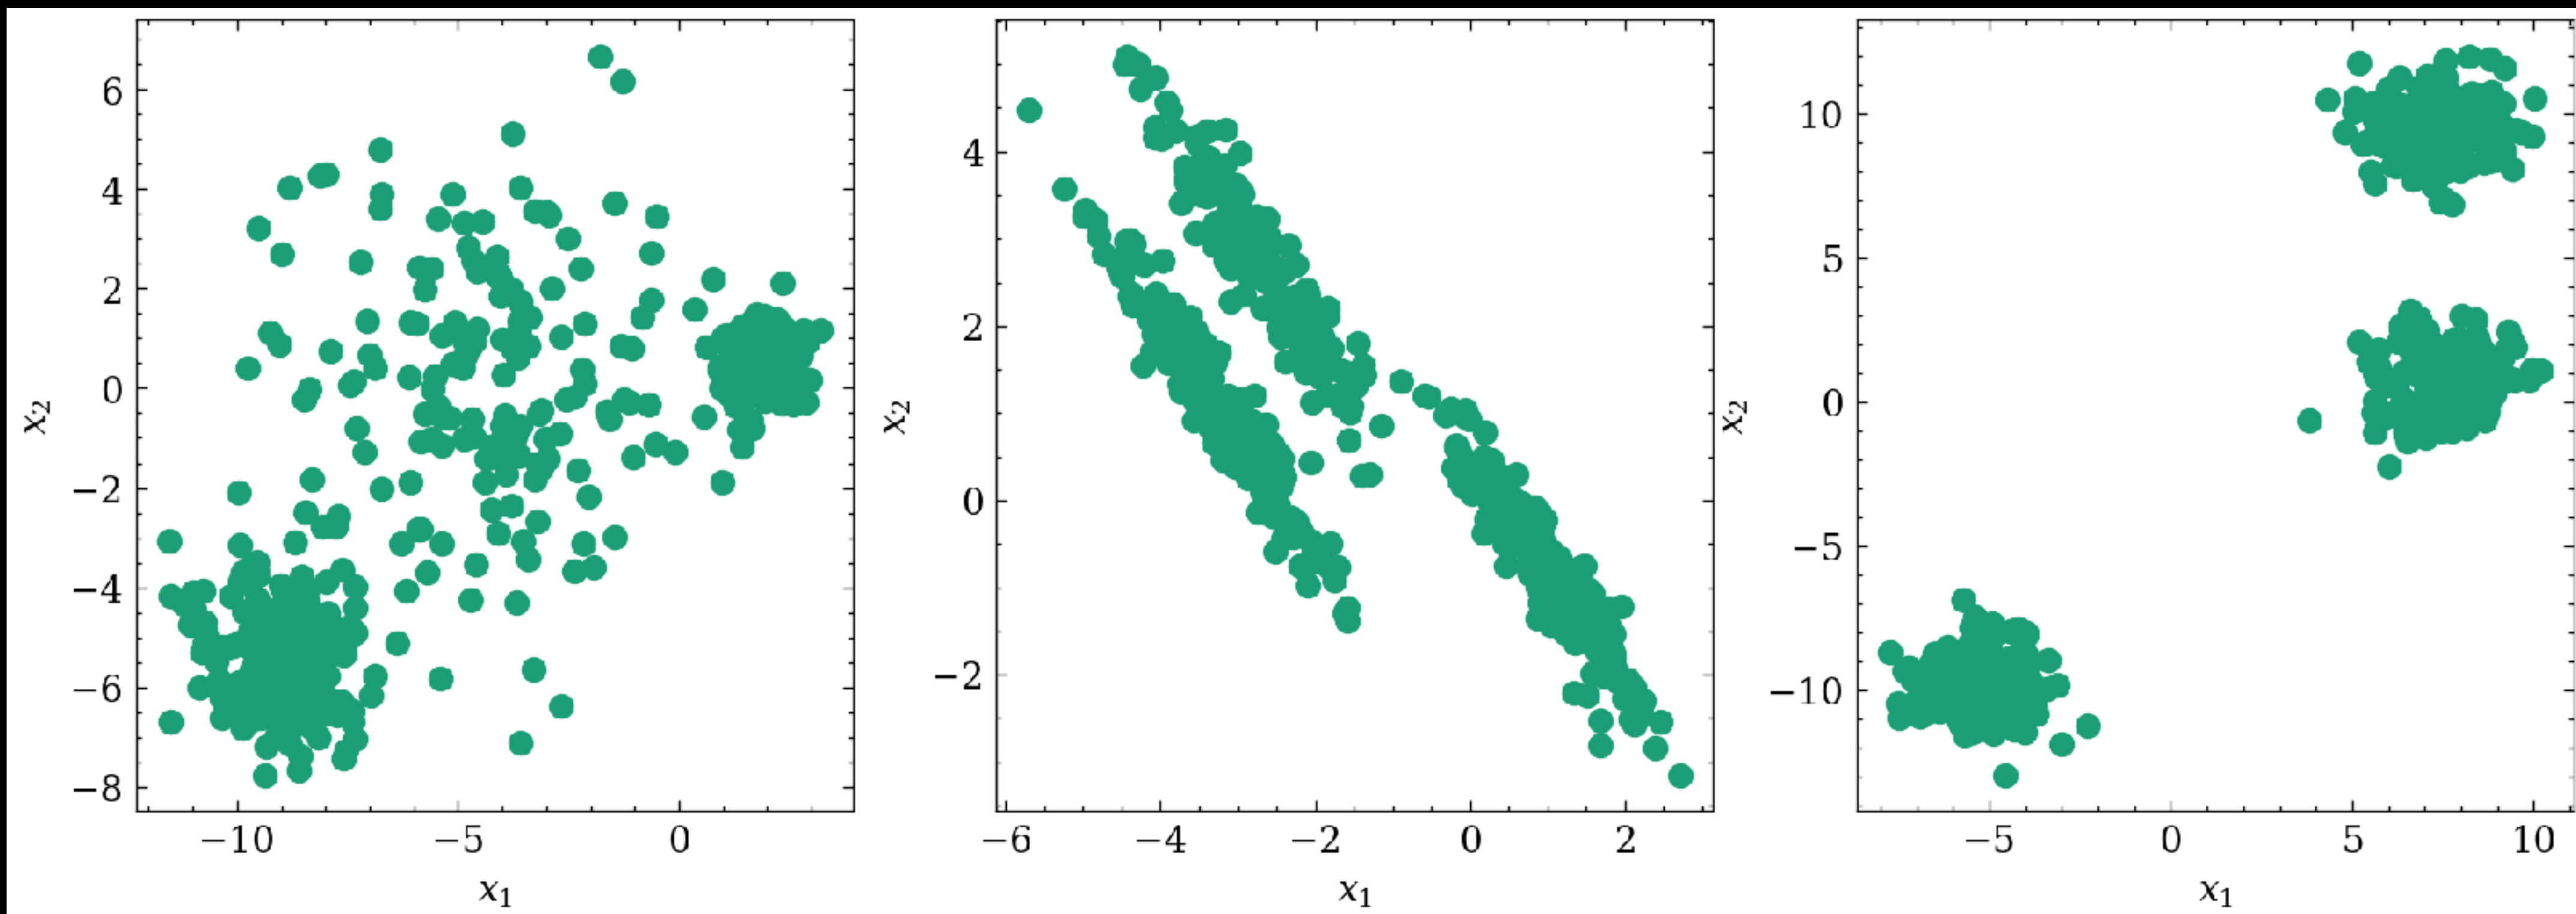
# Clustering

input:

- $x_1, x_2$

Goal: identify M clusters in the data:

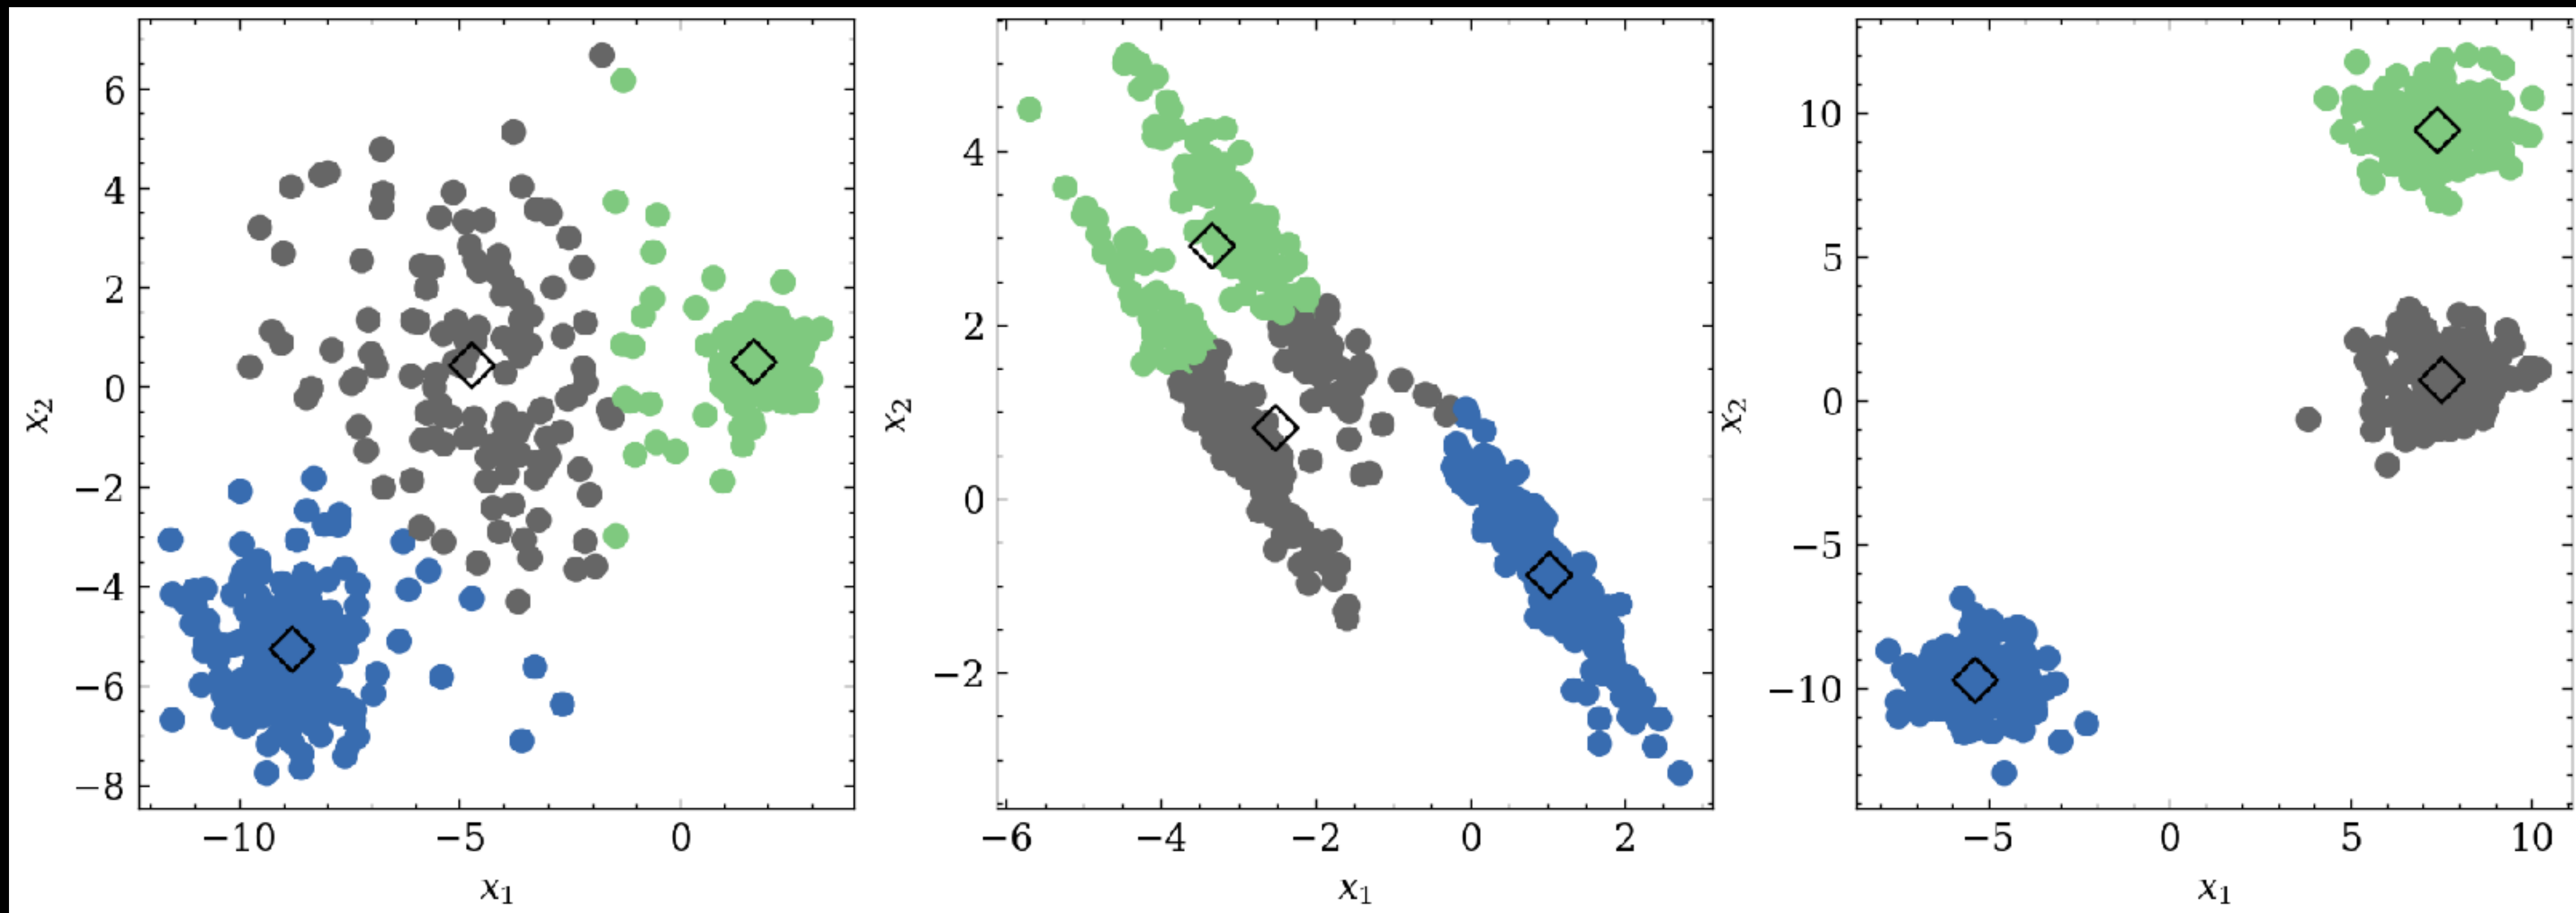
$$\text{predicted class label} = h(x_1, x_2)$$



# K-means Clustering

- requires the number of clusters to be specified
- separate samples in M clusters of equal variance
- choose centroids that minimize the inertia or within-cluster sum-of-squares

$$\sum_{i=0}^N \min_{\mu_j \in C} (x_i - \mu_j)^2$$

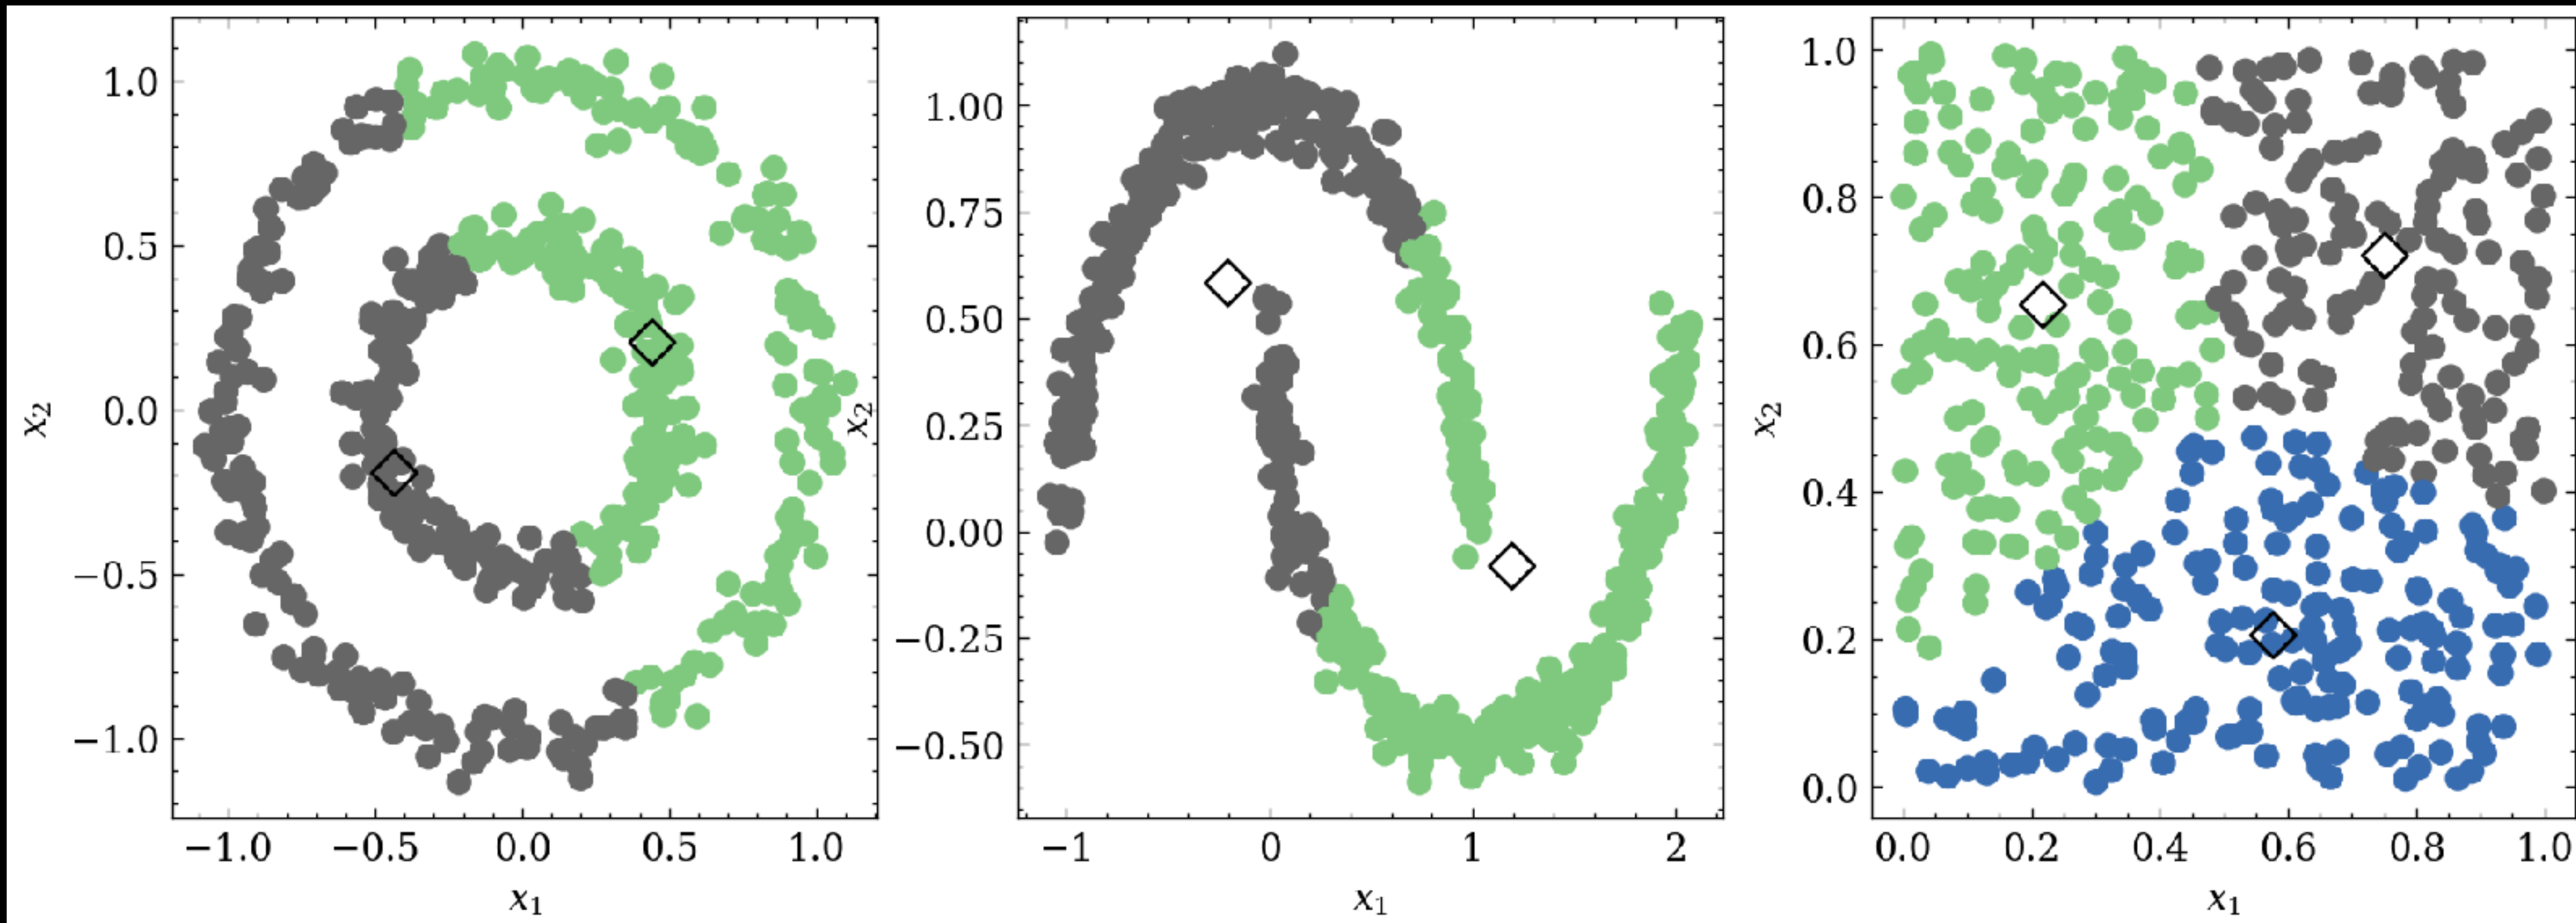




# K-means Clustering

- requires the number of clusters to be specified
- separate samples in M clusters of equal variance
- choose centroids that minimize the inertia or within-cluster sum-of-squares

$$\sum_{i=0}^N \min_{\mu_j \in C} (x_i - \mu_j)^2$$



Examples of problems in your area of research  
where *clustering* could be used?

# Common algorithms

Supervised	Unsupervised
Regression	Clustering
Classification	Dimensionality Reduction

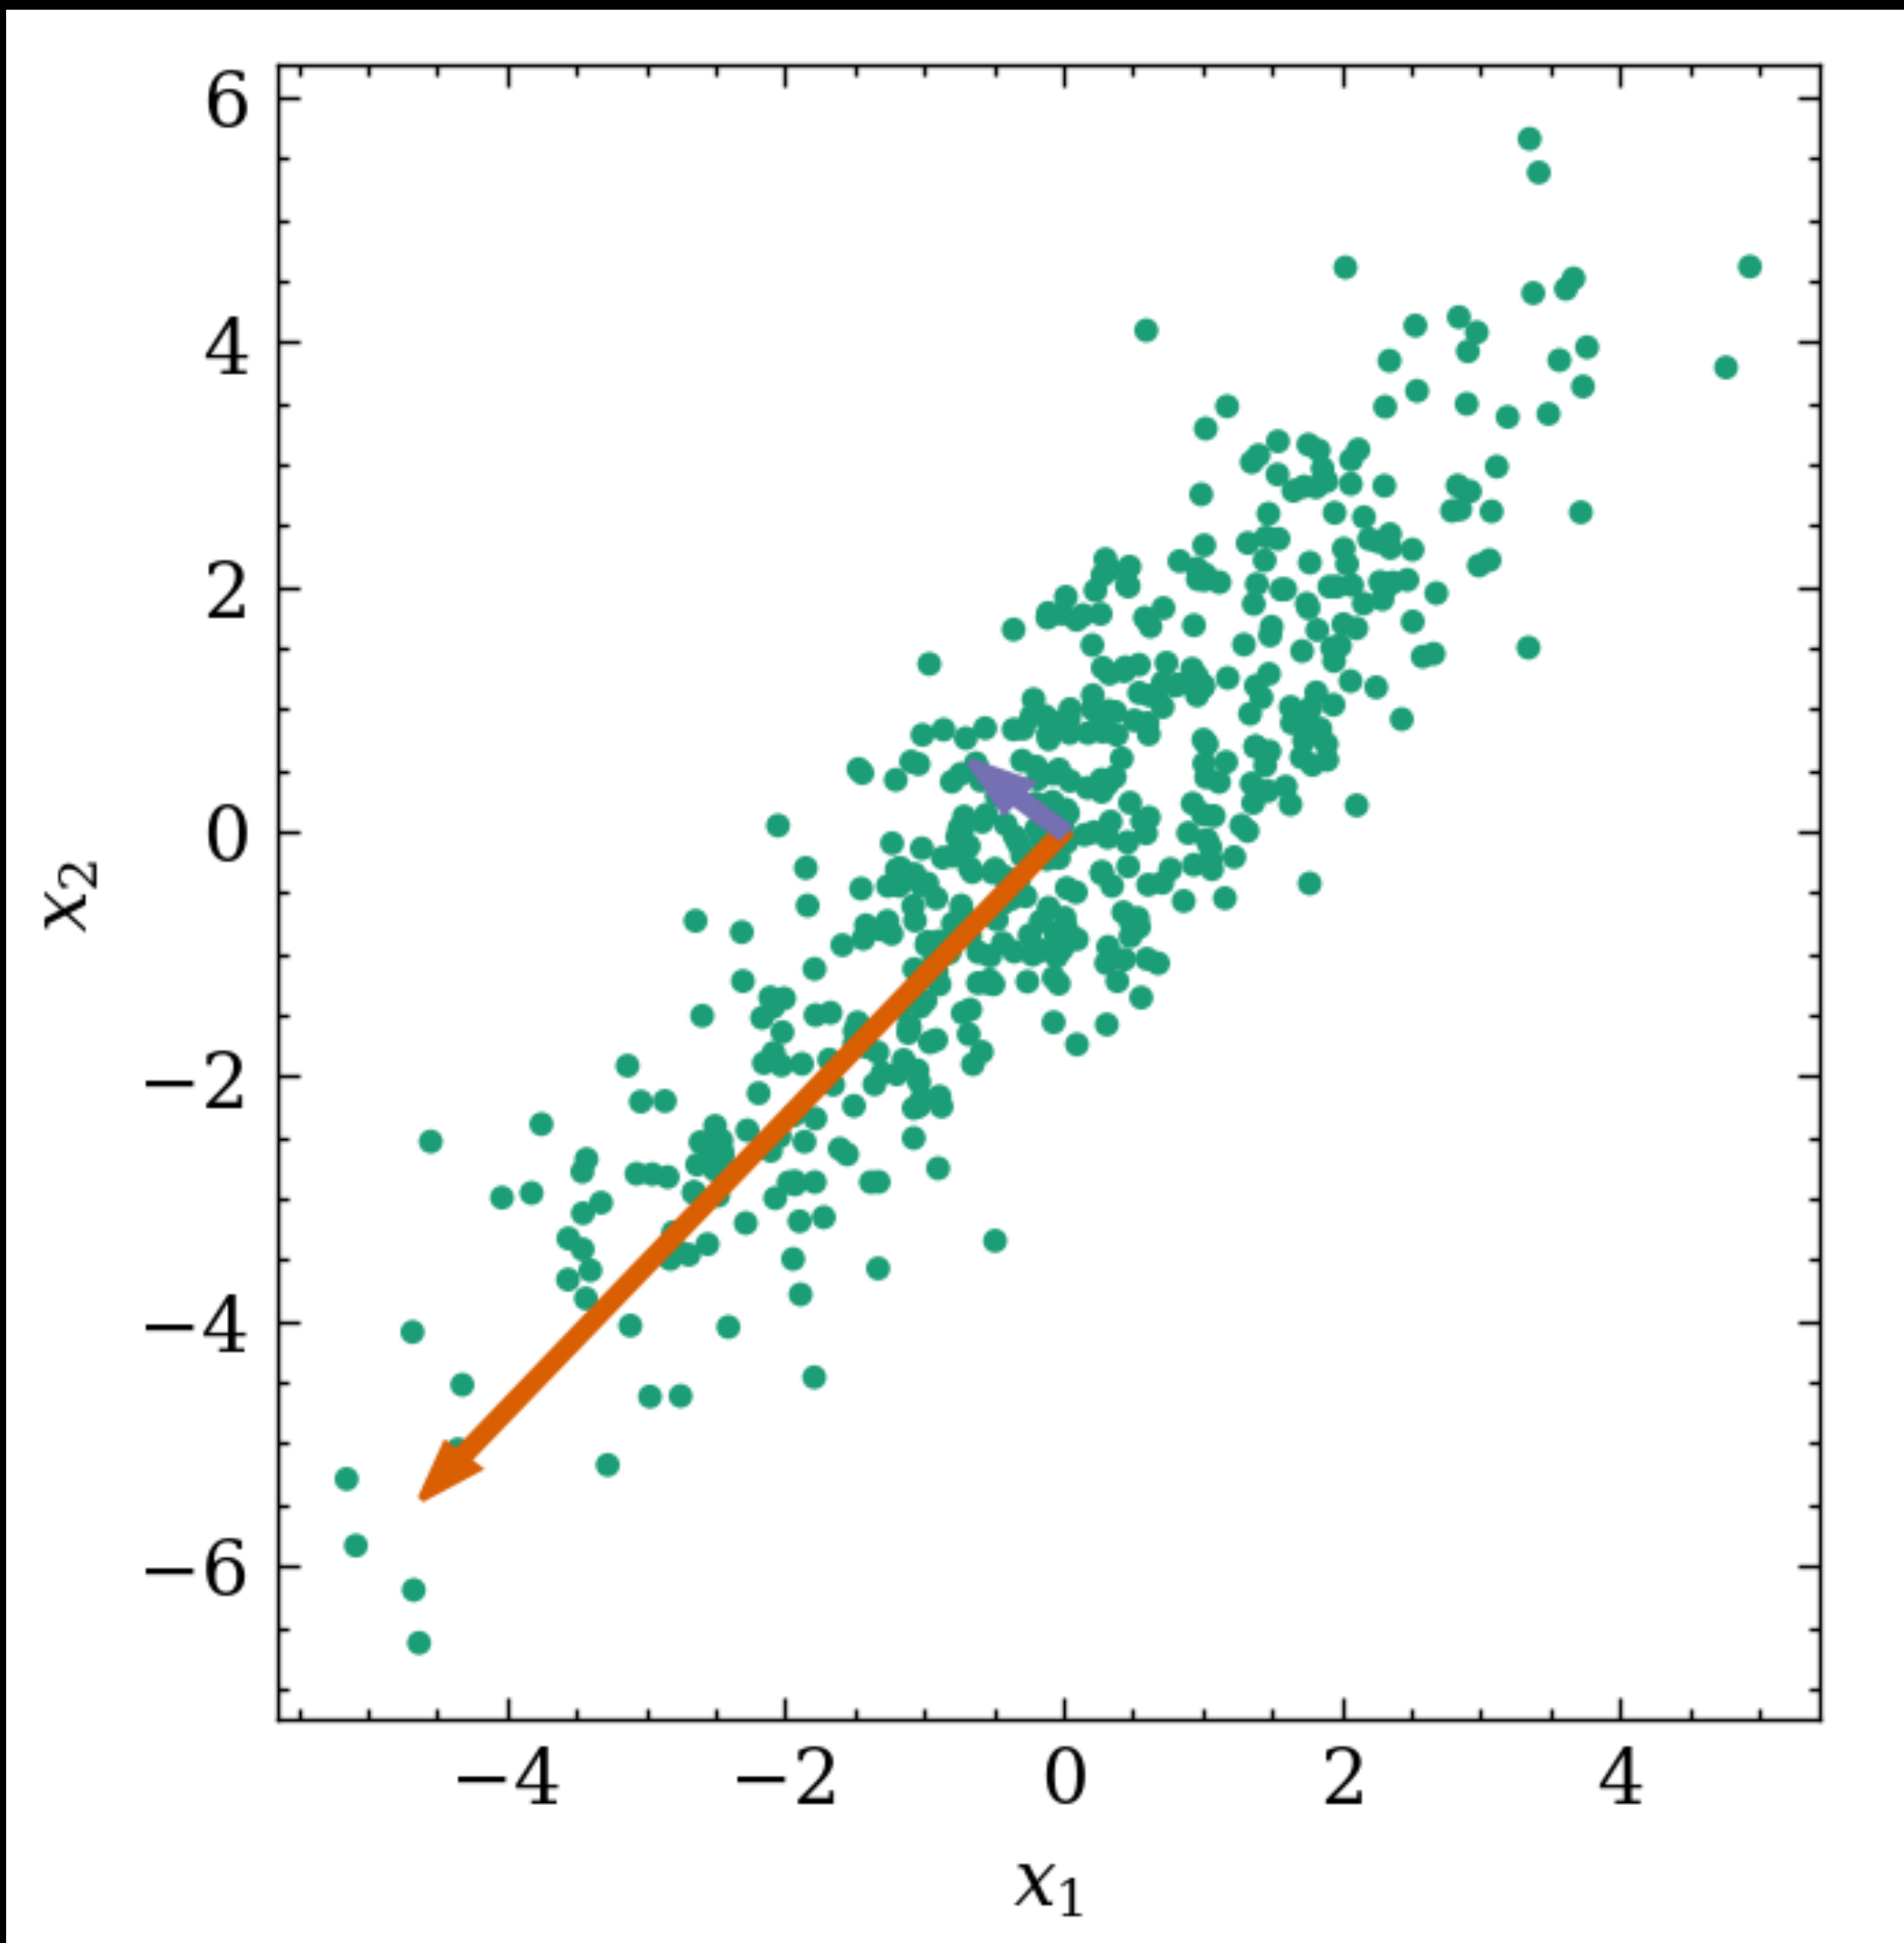
# Dimensionality Reduction

Training set of N inputs pairs:

$X_1, X_2, \dots, X_N.$

Goal: data exploration / visualization

# Principal Component Analysis

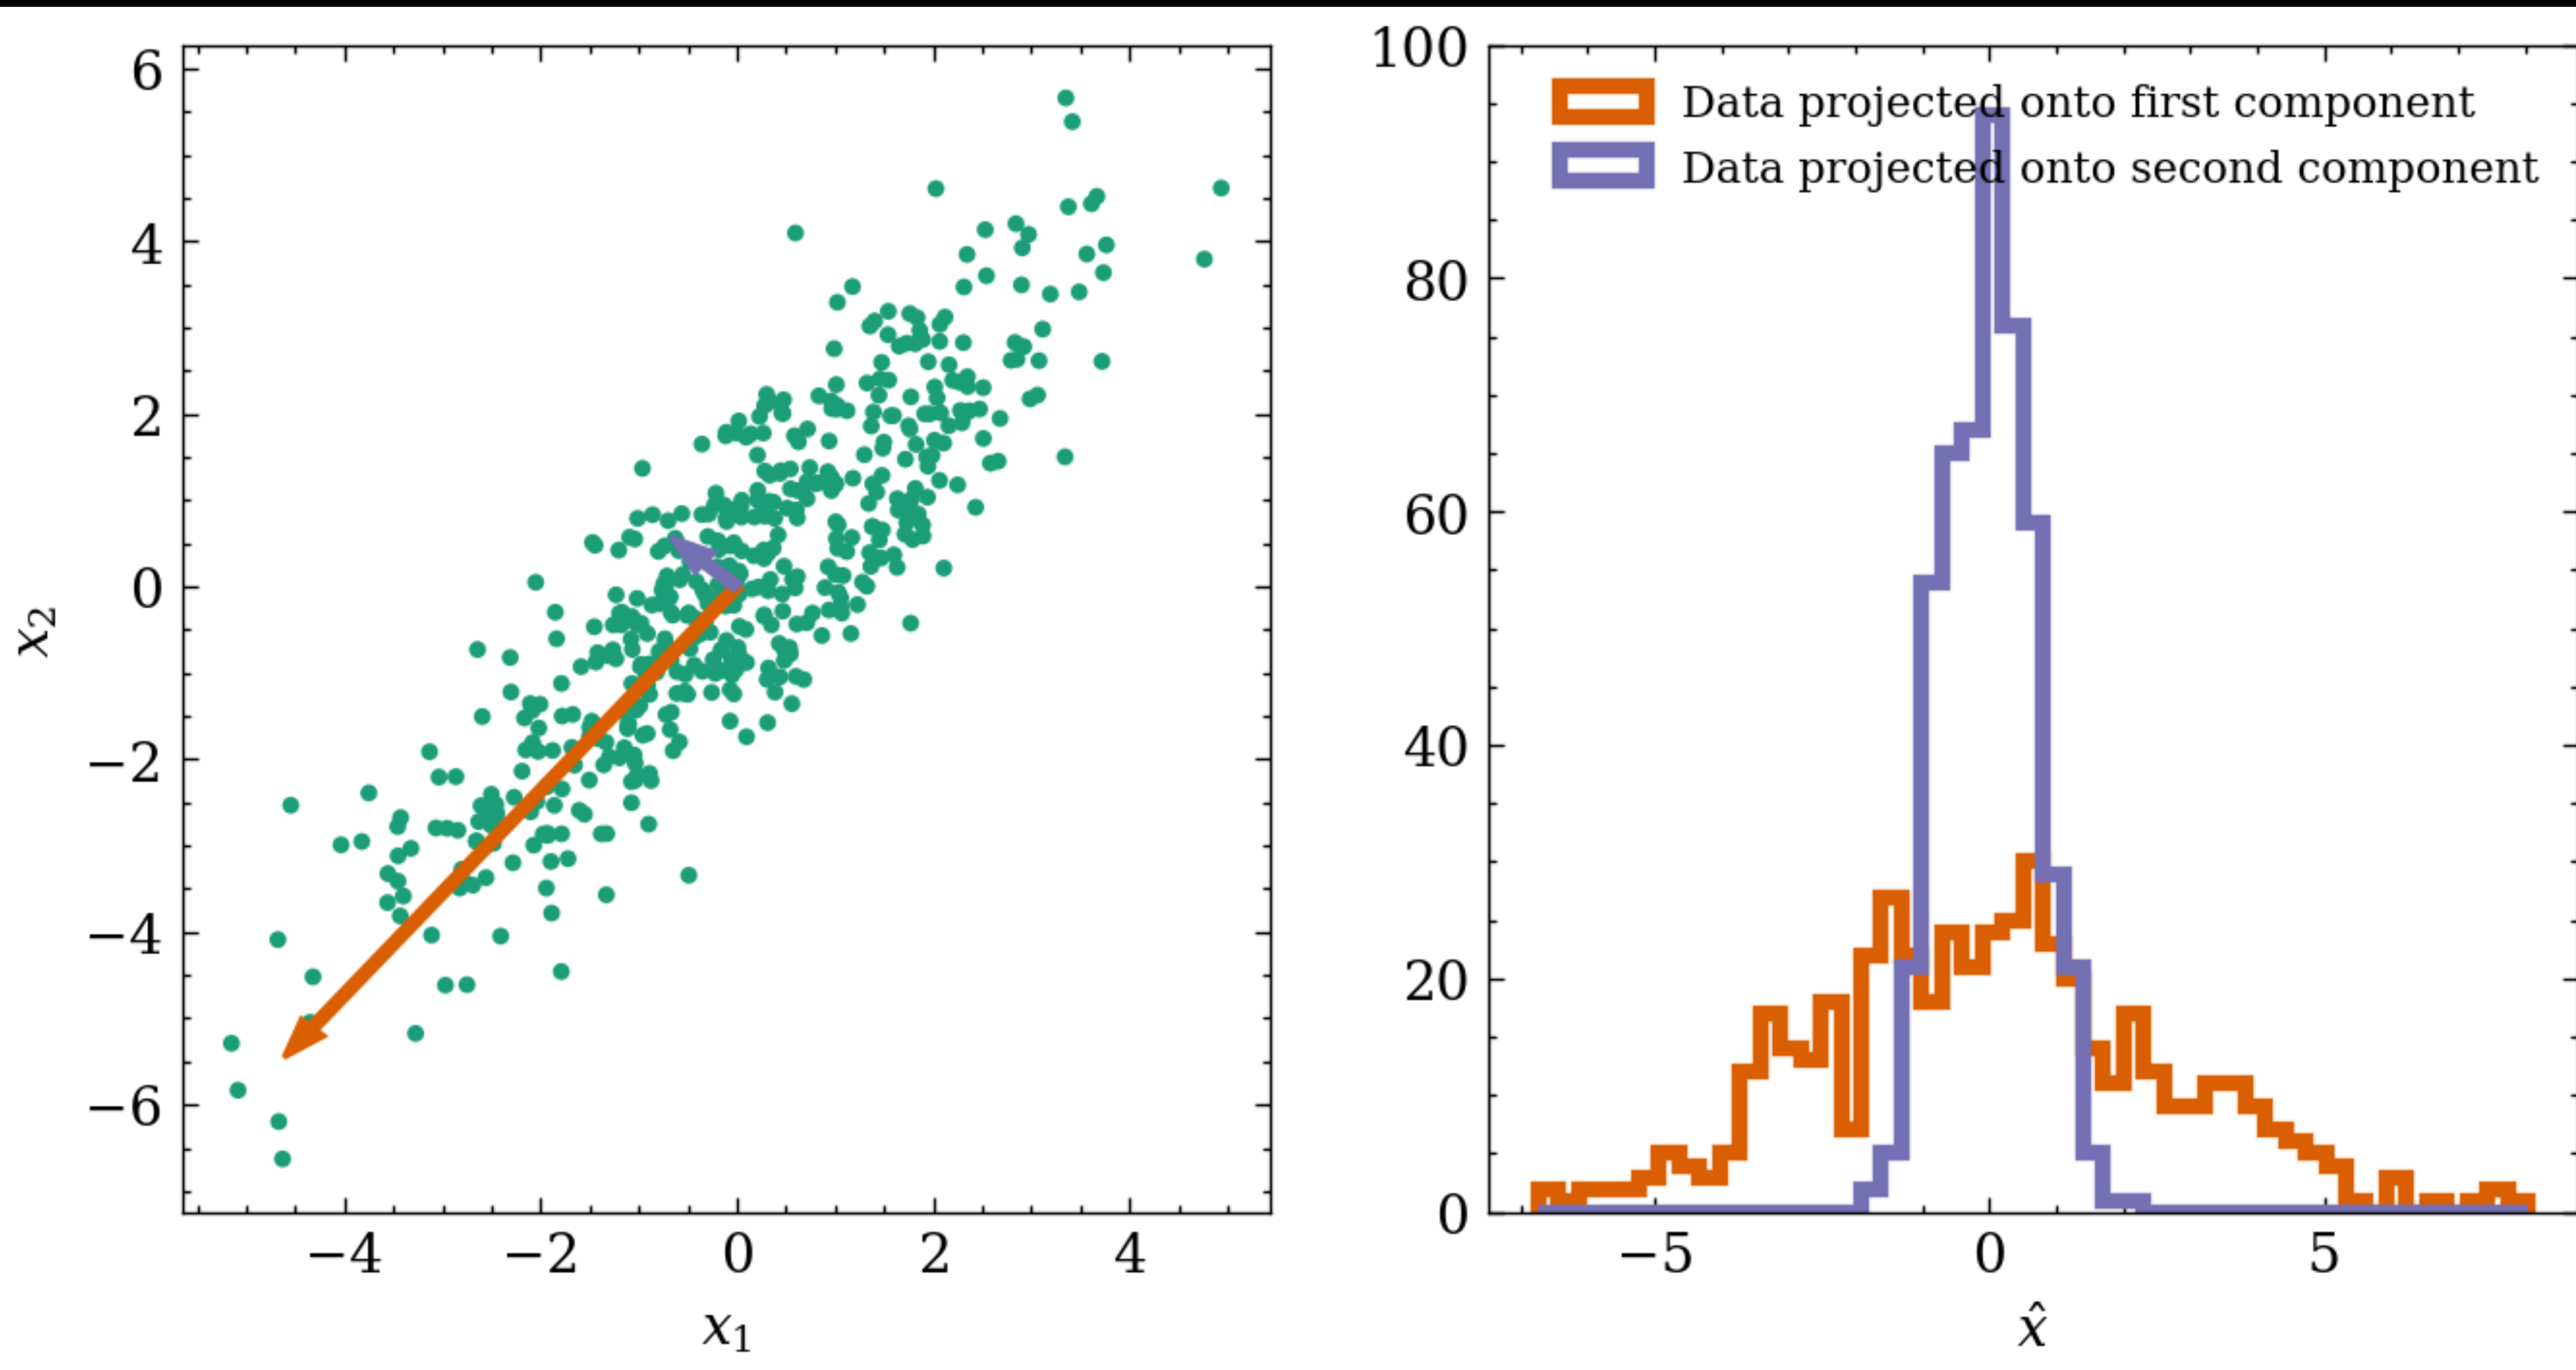


- First principal component: line that minimizes the average squared perpendicular distance from the points to the line.
- Second principal component: line orthogonal to first principal component, that does the same.
- Use the principal components to perform a change of coordinate system.

*sklearn.decomposition.PCA*

*Pearson, K. (1901). "On Lines and Planes of Closest Fit to Systems of Points in Space". Philosophical Magazine. 2 (11): 559–572.*

# Principal Component Analysis





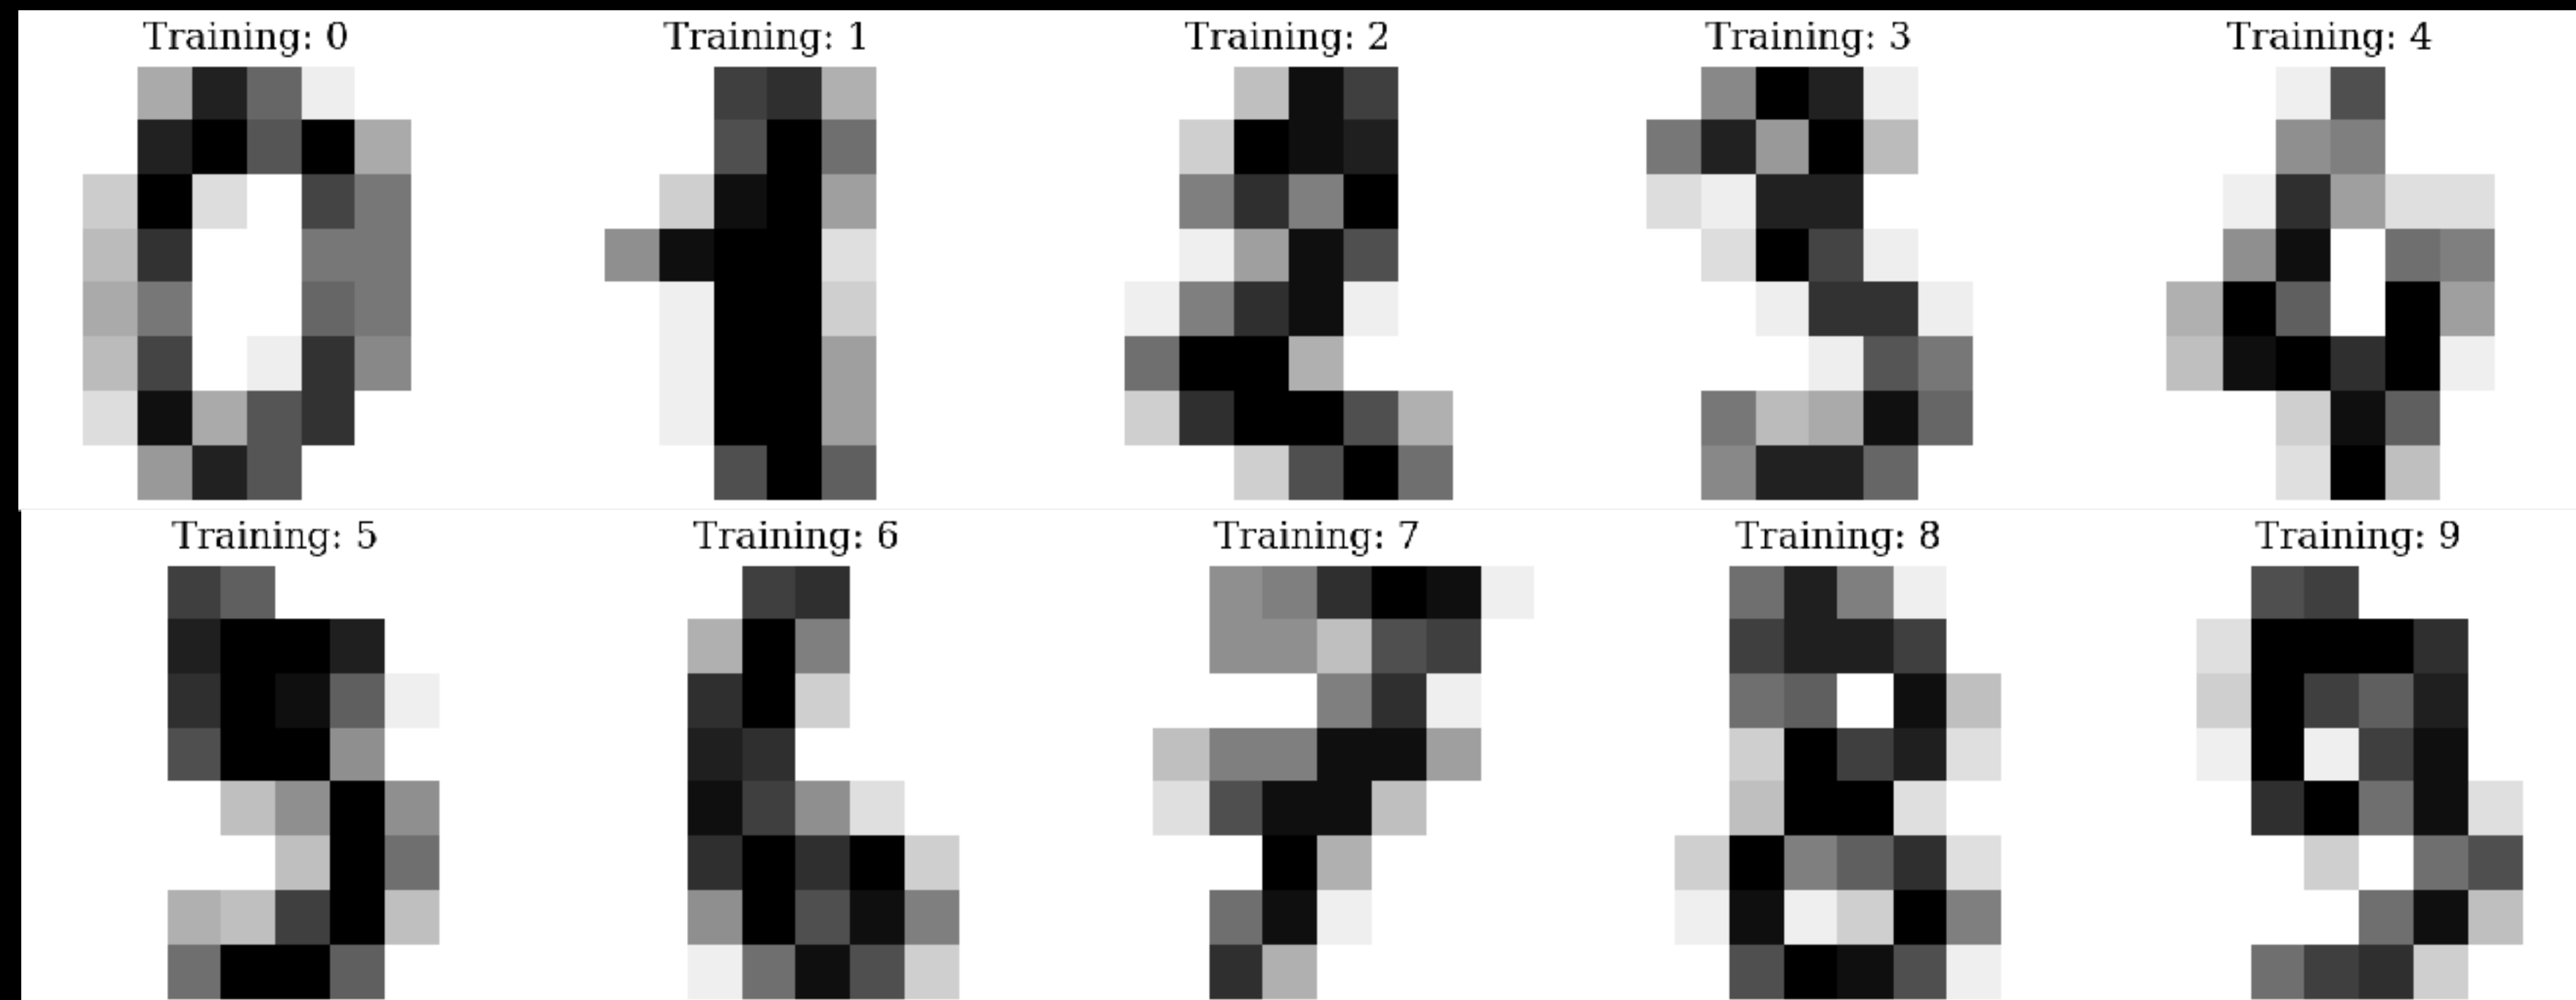
# Dimensionality Reduction

input:

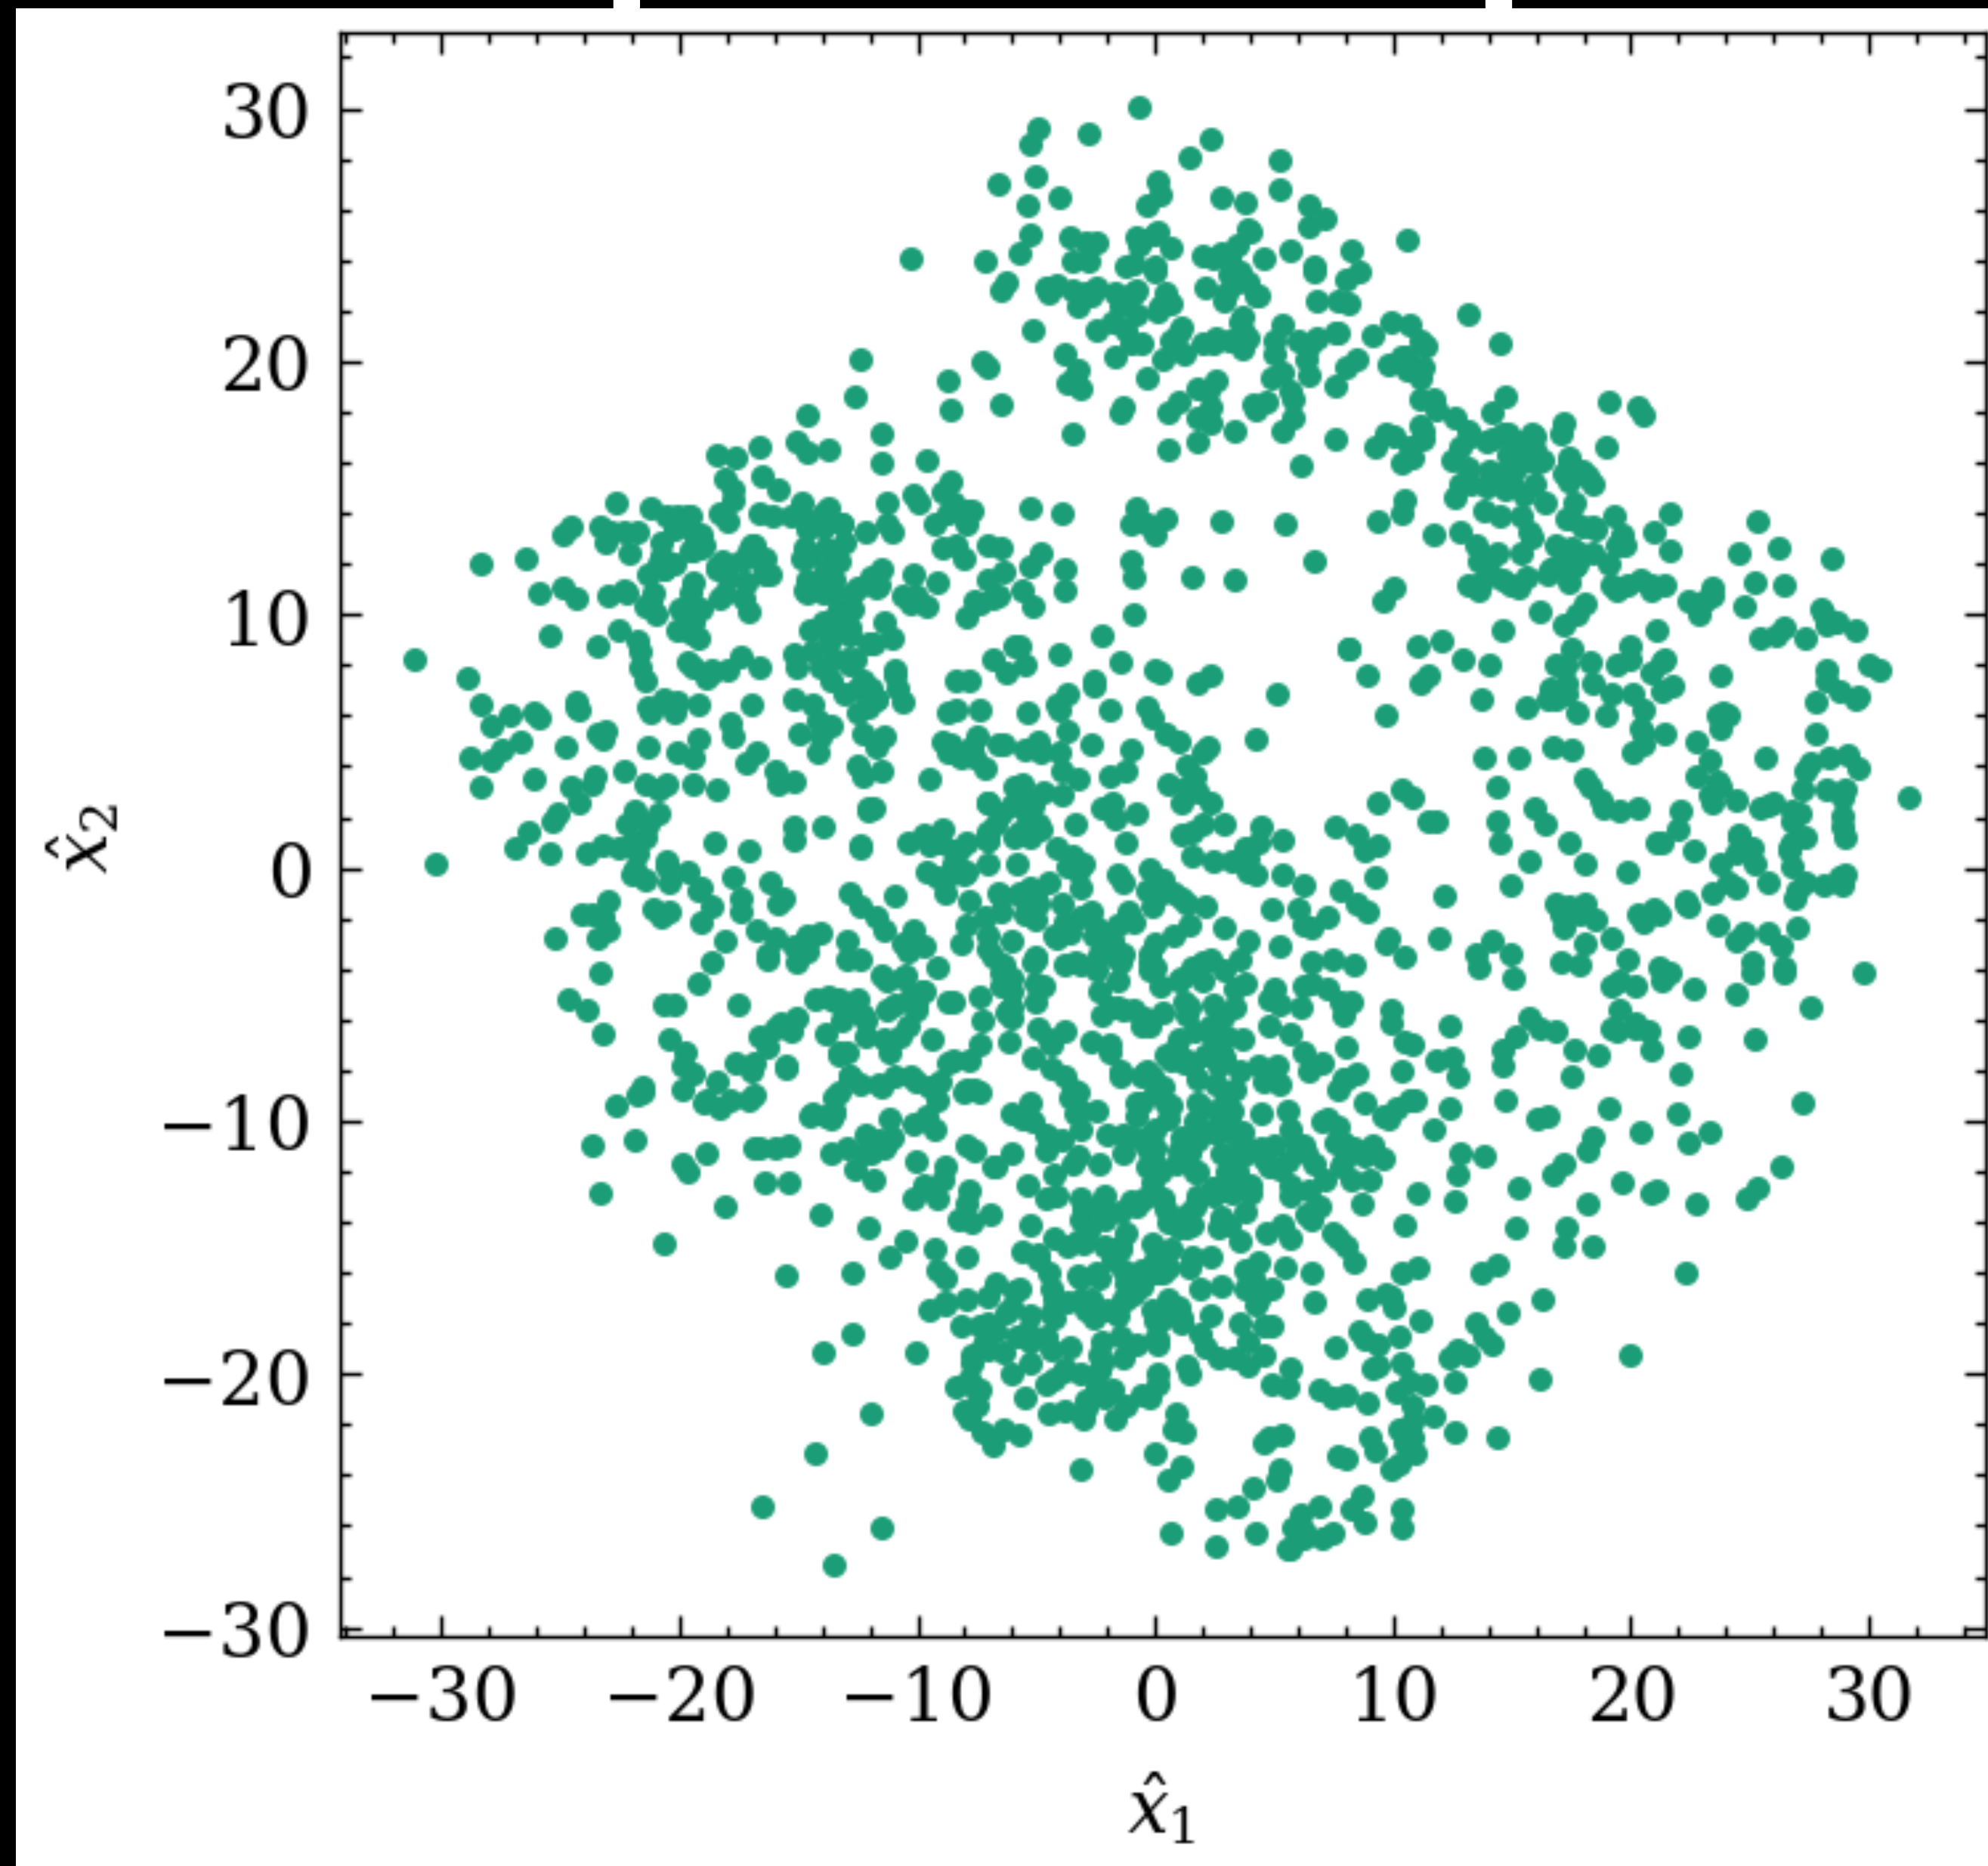
- $x_1, x_2, \dots, x_{64}$

Goal: visualize the data in two dimensions

$$\hat{x}_1, \hat{x}_2 = h(x_1, x_2, \dots, x_{64})$$



# Principal Component Analysis

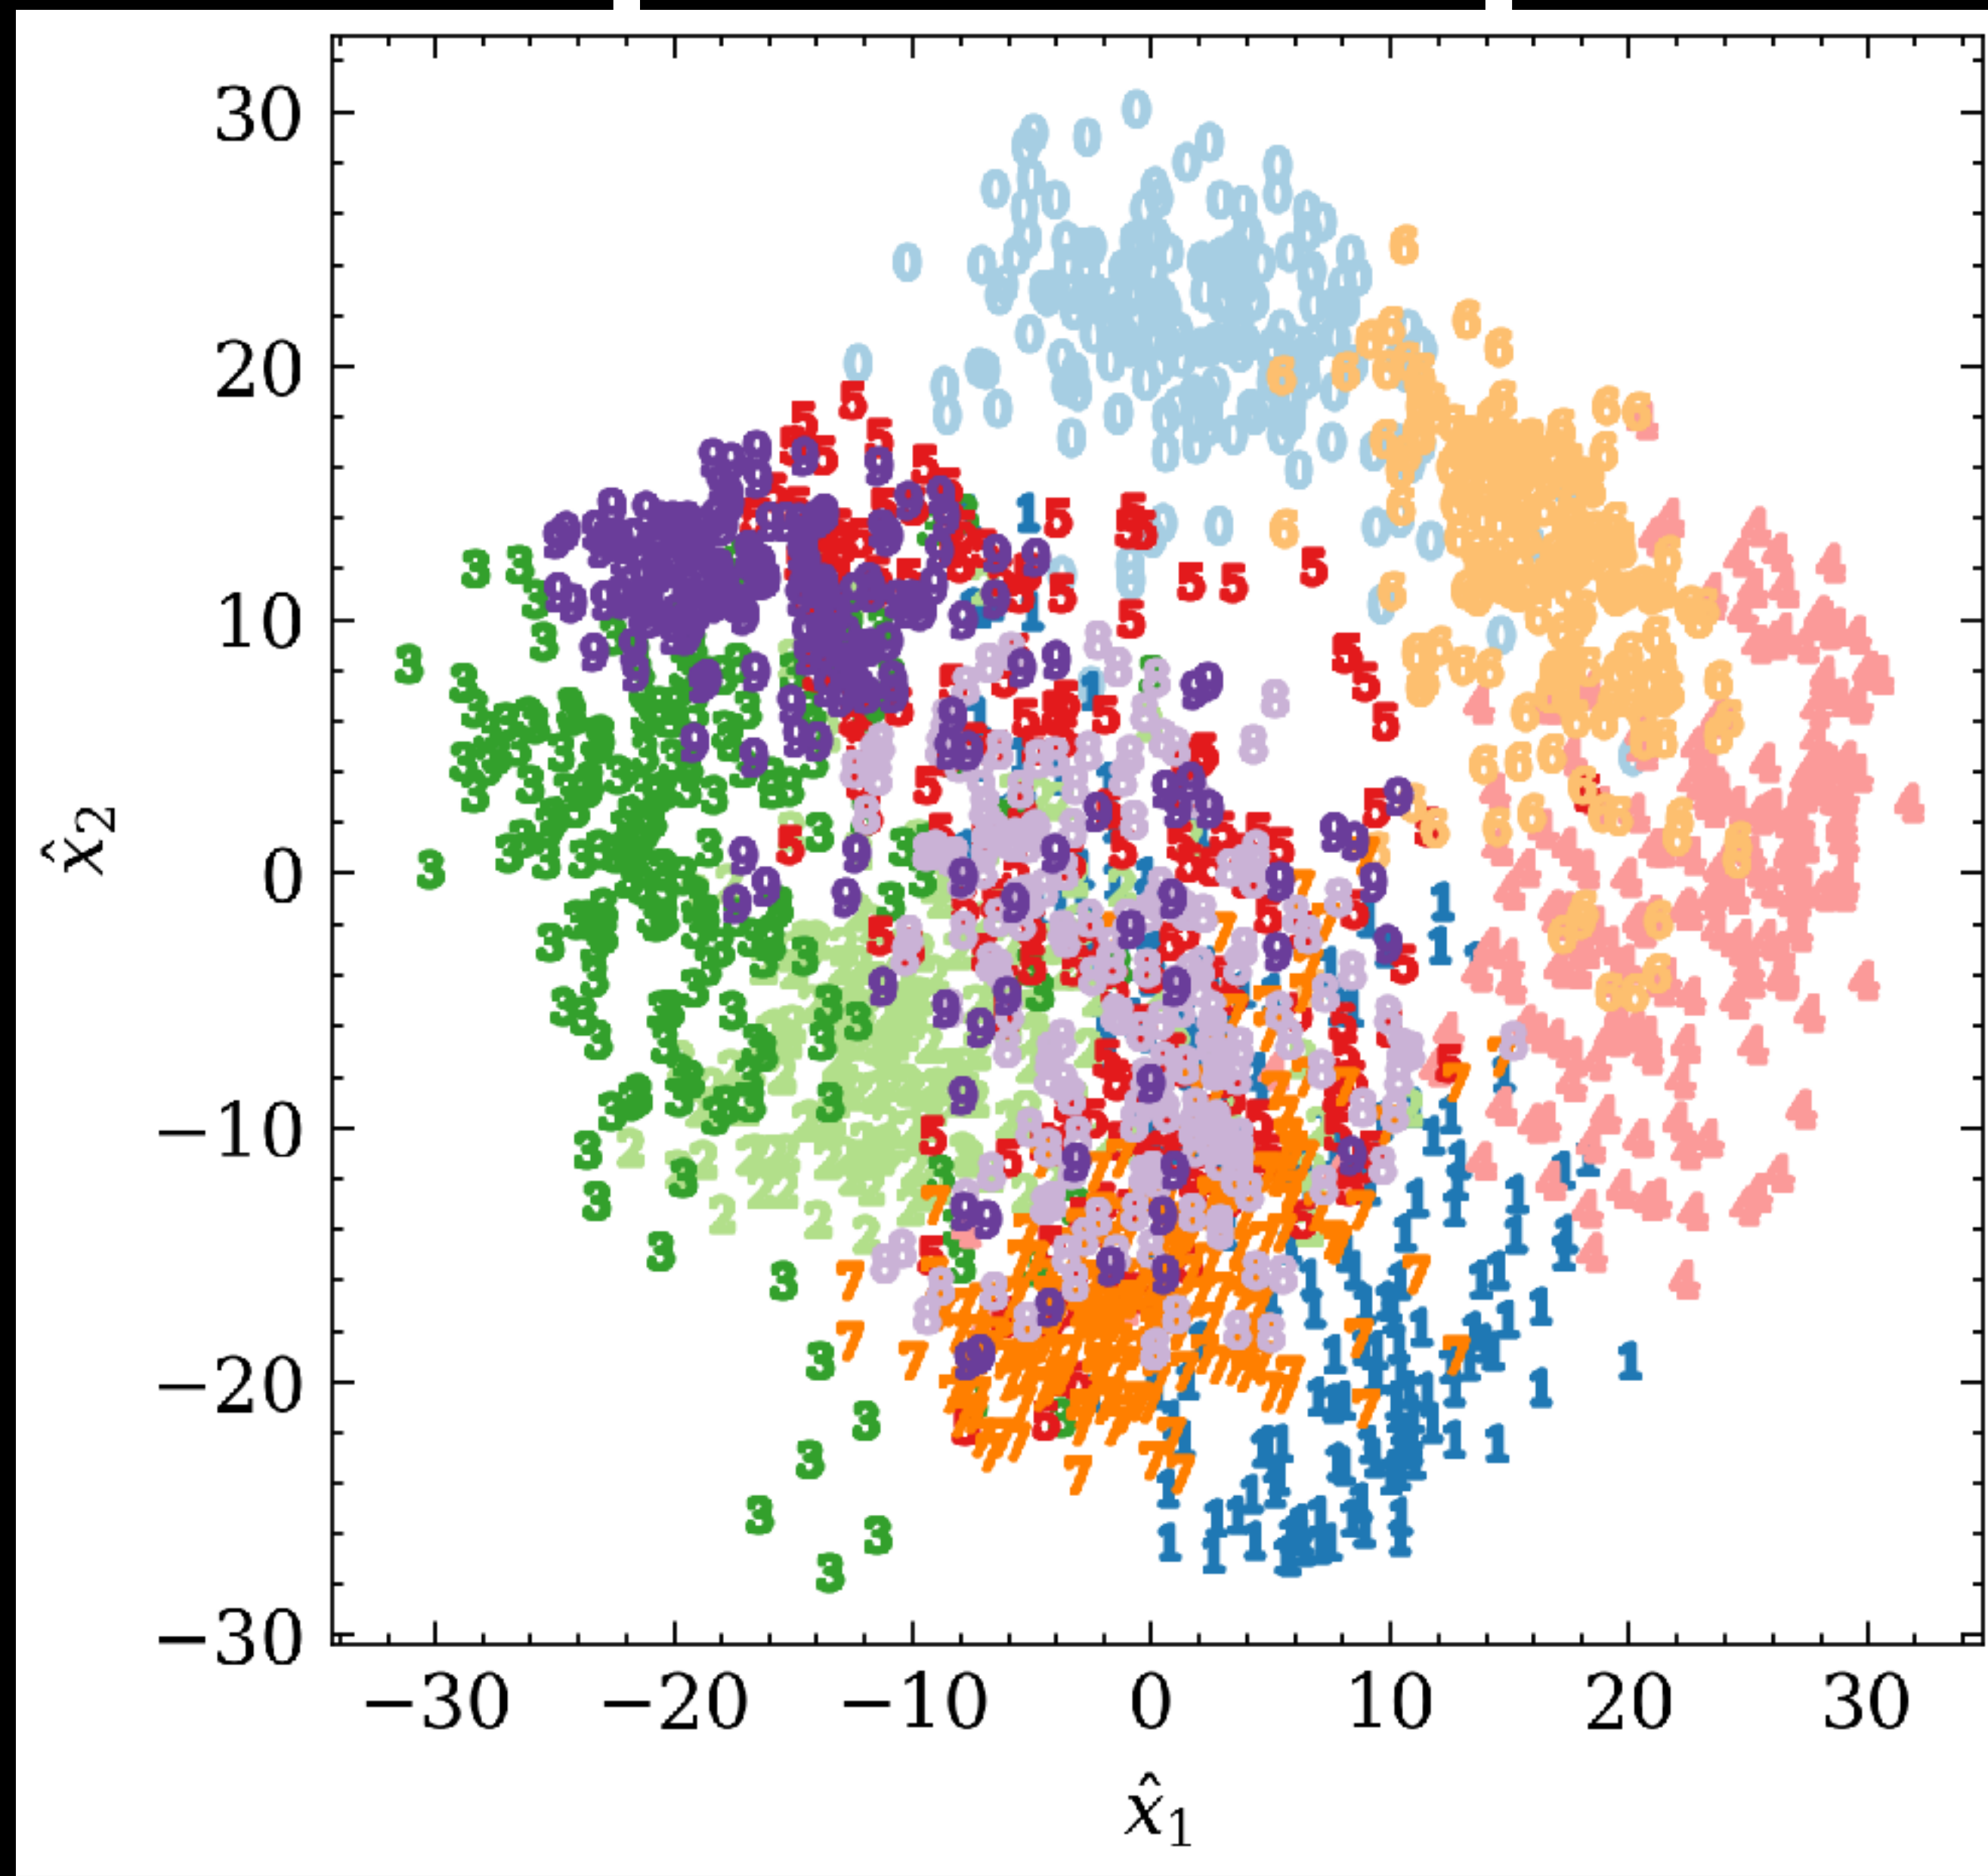


The data set projected down onto the first two principal components

*sklearn.decomposition.PCA*



# Principal Component Analysis



The data set projected down onto the first two principal components, with labels.

*sklearn.decomposition.PCA*

Examples of problems in your area of research  
where *dimensionality reduction* could be used?